GLOBAL PROGRAM OPTIMIZATIONS.

GESCHKE, CHARLES MATTHEW

ProQuest Dissertations and Theses; 1973; ProQuest Dissertations & Theses Global

INFORMATION TO USERS

This dissertation was produced from a microfilm copy of the original document. While the most advanced technological means to photograph and reproduce this document have been used, the quality is heavily dependent upon the quality of the original submitted.

The following explanation of techniques is provided to help you understand markings or patterns which may appear on this reproduction.

- The sign or "target" for pages apparently lacking from the document photographed is "Missing Page(s)". If it was possible to obtain the missing page(s) or section, they are spliced into the film along with adjacent pages. This may have necessitated cutting thru an image and duplicating adjacent pages to insure you complete continuity.
- When an image on the film is obliterated with a large round black mark, it is an indication that the photographer suspected that the copy may have moved during exposure and thus cause a blurred image. You will find a good image of the page in the adjacent frame.
- 3. When a map, drawing or chart, etc., was part of the material being photographed the photographer followed a definite method in "sectioning" the material. It is customary to begin photoing at the upper left hand corner of a large sheet and to continue photoing from left to right in equal sections with a small overlap. If necessary, sectioning is continued again beginning below the first row and continuing on until complete.
- 4. The majority of users indicate that the textual content is of greatest value, however, a somewhat higher quality reproduction could be made from "photographs" if essential to the understanding of the dissertation. Silver prints of "photographs" may be ordered at additional charge by writing the Order Department, giving the catalog number, title, author and specific pages you wish reproduced.

University Microfilms

300 North Zeeb Road Ann Arbor, Michigan 48106 A Xerox Education Company

73-9133

GESCHKE, Charles Matthew, 1939-GLOBAL PROCRAM OPTIMIZATIONS. [Pages 135-137, previously copyrighted material not microfilmed at request of author. Available for consultation at the Carnegie-Mellon University Library].

Carnegie-Mellon University, Ph.D., 1973 Computer Science

University Microfilms, A XEROX Company, Ann Arbor, Michigan

GLOBAL PROGRAM OPTIMIZATIONS

Charles Matthew Geschke

Department of Computer Science Carnegie-Mellon University Pittsburgh, Pennsylvania 15213 October, 1972

Submitted to Carnegie-Mellon University in partial fulfillment of the requirements for the degree of Doctor of Philosophy.

This work was supported by the Advanced Research Projects Agency of the Office of the Secretary of defense (F44620-70-C-0107) and is monitored by the Air Force Office of Scientific Research. This document has been approved for public release and sale; its distribution is unlimited.

Carnegie-Mellon University

CARNEGIE INSTITUTE OF TECHNOLOGY AND

MELLON INSTITUTE OF SCIENCE

THESIS

SUBMITTED IN PARTIAL FULFILLMENT OF THE REQUIREMENTS
FOR THE DEGREE OF Doctor of Philosophy

TITLE	Global Program Optimizations	
PRESENTED BY	Charles M. Geschke	
ACCEPTED BY	THE DEPARTMENT OF Computer Science	Oct 3, 1972 Oct 3, 1972 DATE
APPROVED BY	THE COLLEGE COUNCIL	
_	Commind	Oct ~ 4,197

PLEASE NOTE:

Some pages may have

indistinct print.

Filmed as received.

University Microfilms, A Xerox Education Company

ABSTRACT

The dissertation investigates the optimization of object code produced by compilers of higher level languages. Its primary goal is the isolation of a set of primitives which lead to a concise description and correspondingly concise implementation of program optimizations. In addition to being powerful enough to provide a concise representation, the primitives are also basic enough to apply to a wide range of languages and optimization techniques.

The concept of similarity functions is introduced. A set of new optimizations described in terms of the similarity notion is proposed. A translator is described which implements code motion, redundant expression elimination, and new similarity-induced optimizations using the primitives developed in the dissertation. Examples are presented demonstrating the effect of these optimizations.

ACKNOWLEDGEMENTS

My sincere thanks to:

Professor William Wulf, my thesis advisor, who provided constant encouragement, initiated many of the thesis ideas, and oversaw their development;

Professors Fuller, Grayson, Habermann, and Parnas, members of my thesis committee, who helped shape the final form of the dissertation;

my fellow graduate students, especially the Bliss group, Jerry Apperson, Rich Johnsson, Chuck Weinstock, and Dave Wile, who formed an initial, critical audience to many of these ideas;

Peter, Kathleen, and John, who often endured an absent father; and

Nancy, my wife, whose patient understanding and unfailing support sustained

me throughout my graduate education and especially during these last months

while the dissertation was being completed.

TABLE OF CONTENTS

Abstract	ii
Acknowledgements	iii
Table of contents	iv
Chapter I: Introduction	1
A brief history of optimization	2
	7
Thesis outline	1 2 7 7
Initial assumptions	11
A short Bliss primer	••
Chapter II: Optimization primitives	16
Primitive ordering relations	17
Decomposition of ≺-ordering	21
Similarity functions an introduction	25 .
Code motion optimizations	29
Code motion in linear blocks	3Ø
Code motion in forked control	34
Code motion in loops	40
Code motions and leave expressions	45
	47
Strength reduction	57
Redundant expressions	60
Summary	0.0
Chapter III: Similarity functions	63
Origins of the similarity concept	64
Examples of some similarity functions	7.0
Converting expressions to subroutines	78
Partial post-evaluations in forks	81
Wasp-waisting revisited	82
Generating conditional subroutines	83
Generating Conditional Subsequence	84
Similarity and iterative techniques	86
	87
Summary	
Chapter IV: Examples	88
Kate	89
Quadratic formula	98
Gadd-sub	105
Cpoly	107
S3 optimization and execution time	108
Summary	110
Olastas V. Canalysian	113
Chapter V: Conclusion	113
Summary of thesis results	114

Appendix A: Gadd-sub			119
			135
Appendix B: Cpoly			138
Glossary of definitions	•		139
Bibliography			. 138

CHAPTER I

Since the advent of the first FORTRAN compilers, the loss in object code efficiency incurred by the use of higher-level languages has concerned both programmers and compiler designers alike. The proponent of a language intended for compilation, even though he may argue that the cost in lost efficiency is far outweighed by the power and elegance of his language, must generally supply a compiler which produces reasonably efficient code in order to attract a community of users. The new breed of "languages for implementation of systems" is measured against this criterion of efficiency in the extreme.

This thesis investigates the area of object code optimization in the presence of control flow. Its major goal is the isolation of a set of primitives which lead to a concise definition and a correspondingly concise implementation of program optimizations. In addition to being powerful enough to provide a concise representation, these primitives must also be basic enough to apply to a wide range of languages and optimization techniques.

The search for a set of primitives to describe a collection of varied optimizations is motivated initially by a desire to achieve a uniform

representation of these optimization strategies. A uniform representation, in turn, leads to an implementation which can be easily structured into combinations of the set of primitives. As a result the same clarity and concision which is inherent in the primitives is reflected in the implementation. In order to demonstrate this correlation between the description and implementation of various optimizations, a later chapter will discuss the structure of an actual optimization pass within a real compiler which uses the primitives.

The identification of a collection of primitives produces another benefit. The ability to perform formal manipulations on these primitives aids in exposing new optimization strategies and helps identify the common characteristics of apparently unrelated techniques. This effect is, of course, more difficult to document. It has been our experience that even though the discovery of an optimization strategy may not develop solely from manipulating the primitives, the ability to grasp the essential characteristics of an optimization is significantly enhanced by the availability of a set of objects which can be used to describe that strategy concisely.

A BRIEF HISTORY OF OPTIMIZATION

Our investigation has evolved through a set of selections among various alternatives and been motivated by several goals, some already

described above and others yet to be stated. Any such evolution of ideas builds upon the work of our predecessors who have investigated the problem of object code optimization. We will not attempt to produce a complete catalog but instead will select those efforts which have guided our choices among alternatives either by contrast or in parallel.

In June, 1965 an article by J. Nievergelt[N65] provided a principle for this area of investigation that seems to remain valid today. He states of optimization extent the constraint on limiting corresponding to our own: a programmer can optimize his program by relying to a great extent on his knowledge of what that program is to do. Indeed his initial encoding of the solution was already a significant optimization some less well-defined general problem solving technique. optimizations we consider are restricted to those which depend on the form of the program only. The results of this thesis show that there continues to be a significant gain in object code efficiency resulting from this level of optimization. As the sophistication of programming languages progresses, it becomes the responsibility of the optimizing compiler to remove the burden of the more tedious details of low-level optimizations from the user. Indeed, as the class of operators and the complexity of data structures grow in power and breadth, the programmer becomes further removed from the target machine (as does the language, perhaps). At some point, then, he is no longer capable of dealing with (or better; he should no longer be as concerned with) the complexities of optimizing his

constructs.

In August, 1965 C.W. Gear[G65] summarized and collected information on the state-of-the-art of machine independent optimizations and proposed a three pass compilation incorporating those strategies. That collection of optimizations remains the basis for most of today's investigations.

A significant amount of research into the area of optimization has centered around the work of F. Allen[A69,A70], J. Cocke[C70], and J. Schwartz[CS70]. Their influence is very evident in the optimizations of the FORTRAN-H compiler which are described by E. Lowry and C. Medlock[LM69]. The authors state that at the cost of a 40 percent increase in compilation time they produce code which is 25 percent smaller and which executes in one-third the time of that produced by the FORTRAN-G compiler. These measurements indicate the real effectiveness of the collection of optimizations implemented in FORTRAN-H.

Much of the work done by Allen and Cocke concerns itself with the processing of the control flow structure of programs and hence contains a considerable amount of graph-theoretic investigation related to control flow representation. We have chosen instead to restrict the control flow semantics to a go-to-less form of control as exhibited in Bliss[B71,WRH71] and concentrate on primitives which relate to the data flow semantics of a program. These data flow primitives concentrate heavily on exposing the issue of re-ordering evaluations in a language independent manner. Since

the suggestion to eliminate the goto by Dijkstra[D68], a debate has proceeded on the nerits of the proposal[H72,W71,W72]. Our own experience in reading, writing, and compiling go-to-less programs (in the Bliss sense) supports the adoption of this programming style. Moreover, the assumption of this form of control flow has had a significant impact on our investigation of optimization since it enables us to enumerate a small set of control environments and restrict our attention to optimizations related to those control structures. Previous investigations into optimization techniques described in the more general control flow environment, in general, assume that the program can be converted to a representation which is essentially modeled by the control flow semantics of Bliss.

The preliminary notes written by Cocke and Schwartz[CS70] appear to be the single most comprehensive catalog of optimization techniques available. Throughout the thesis we will refer to the collection of optimizations described in that text as the set of "classical" optimization strategies. The text by Cocke and Schwartz provides us with another motivation for proposing a set of primitives. Most of the descriptions of optimization techniques and their implementations are presented in terms of algorithms which often cover several pages and which are closely related to intermediate representations of the program. A major point in introducing our primitives is to demonstrate an alternative method for describing and implementing optimizations which is considerably more concise, understandable, and independent of the intermediate representation.

Finally, anyone investigating the area of optimization must be aware of the interaction of this area with the study of the equivalence of programs and the detection of potential parallelism in a computation. The issue of equivalence of programs arises from recognizing that an optimization strategy is concerned with transforming a program P to a program P' which is input-output equivalent to P. The area of program equivalence is broad in scope but there has been some work done by A. Aho and J. Ullman[ASU70,AU70] from the viewpoint of an application to optimization. In general, however, their work has been restricted to straight-line programs.

Many optimization techniques involve the re-ordering of the evaluation of expressions in a program. Equivalently those expressions, whose order of evaluation can be interchanged, can in fact be executed in parallel with sufficient interlocks. Some very interesting work in representing the inherent parallelism in a program has been done by R. Shapiro and H. Saint[SS69] using Petri Nets. While the Petri Net model provides an elegant framework for their investigations, this thesis proposes primitives which are more easily implementable in the environment of a compiler.

In addition to the influence of the above work, another principle has directed our selection among several areas of program optimization. We intend to investigate only machine independent optimizations. Thus, for example, we will not discuss "peephole" otpimization. Typically optimizations of this class exploit the instruction set of a particular

computer by combining a sequence of several operations into a single machine instruction. Also the thesis will not investigate the area of register allocation. Although this area still requires extensive investigation, the time space constraints on a dissertation have led us to concentrate on those machine independent optimizations which most directly evolve into the new optimizations presented later in the thesis.

THESIS OUTLINE

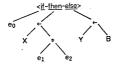
The thesis contains five chapters and two appendices. The remainder of the introduction summarizes our initial assumptions and gives a brief introduction to Bliss. Chapter II introduces the primitives and describes various optimizations techniques in terms of those primitives. Chapter III discusses a concept called similarity which is then used to describe an additional collection of new optimization techniques. Chapter IV presents a set of examples illustrating the various optimization strategies proposed in Chapters II and III. Chapter V contains a summary of our results and suggestions for future research.

INITIAL ASSUMPTIONS

It is inappropriate that a thesis in the area of optimization should tie itself to a single language or single target machine. On the other

hand some assumptions are necessary to form the starting point of an investigation. The viewpoint adopted throughout the thesis holds that the optimization algorithms operate on a tree representation of the source program. The syntax analyzer produces a tree in which each control environment and each operator of the source program is represented by a unique node. In the case of an operator its subnodes are its operands whereas in the case of a control environment the subnodes are its subcomponents. For example the program text

<u>if</u> e₀ <u>then</u> X←e₁*e₂ <u>else</u> Y←B is represented as



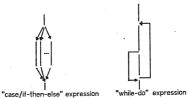
Terminal nodes are always literals or names. The following notational convention is observed for a node, e, such as the <u>if-then-else</u> expression above:

e[operator] = $\langle \underline{if}$ -then-else> e[# of operands] = 3 e[operand₁] = e₀ e[operand₂] = X \leftarrow e₁ \approx e₂ e[operand₃] = Y \leftarrow B

The goal of the optimizer is to produce a transformation of this tree which is more optimal in accordance with whatever time/space guidelines the target machine (and perhaps) the user has imposed. The variability in

target machines is factored out of the optimizations strategies by: (1) allowing input of the characteristics of the target machine to decision making procedures of the optimizer, and (2) requiring that the optimizer encode sufficient information for the code generators and temporary storage allocators of a particular machine to use in their decisions.

The validity of any reshaping of the program tree is dependent upon the semantics of both the control flow and data flow of the source language. The discussions involving references to control structures are couched in terms of flow diagrams such as the following.



While they have their obvious counterparts in the syntax of many languages, all development is independent of a particular syntax.

The major assumption about the control flow semantics, as was stated above, is that the language is go-to-less. The thesis does not consider the problem of detecting programs which fit this model nor the problem of transforming programs into this form. This area has been examined extensively by a number of people. As a result of this go-to-less assumption the tree emitted by the syntax analyzer gives a complete

representation of the control flow semantics without further analysis.

In addition to treating the control flow semantics in a general fashion, we wish to factor out of the development the issue of side-effects which result from the semantics of the language's data flow. To this end, a primitive relation, essential predecessor, whose function is to remove the language dependent issue of side-effects, will be introduced. Given a particular language, the semantics of the applications of side-effects within that language define this relation.

The initial assumptions of the thesis are summarized:

- algorithms employing the optimization primitives assume a tree representation of the source program as input and produce a similar representation as output;
- (2) target machine independence is achieved by parameterizing the optimization algorithms and requiring them to produce information for subsequent machine dependent optimizations in the output representation:
- (3) the control flow semantics of the source language are assumed to be go-to-less; and
- (4) language dependent data flow semantics are to be isolated by primitive ordering relations so that subsequent development becomes language independent.

A SHORT BLISS PRIMER

Throughout the thesis we will present examples to clarify and motivate concepts as they are introduced. Bliss (and occasionally Algol[AL60]) will be the languages used in these examples. We emphasize that Bliss is introduced for use as a syntactic representation of the control structures and its use does not reduce the language-independence of the optimizations. Bliss is sufficiently Algol-like in many aspects so that a brief introduction to the language should be sufficient for understanding the examples. More detailed information on Bliss is available elsewhere [871,WRH71].

INTERPRETATION OF NAMES

A Bliss program operates with and on a number of storage "segments". A segment consists of a fixed and finite number of "words". A word may be "named"; the value of a name is called a "pointer" to the word. Identifiers are bound to names by declarations. Thus the value of an instance of an identifier, say x, is not the value of the word named by x, but rather a pointer to x. This interpretation requires a "contents of" operator for which the symbol "." has been chosen.

This context independent interpretation of identifiers as pointers is maintained consistently throughout the language. It is the operators of

Bliss which place an interpretation on the value of an expression. So, for example, the assignment operator "\infty" interprets its right hand operand as a value which is to be stored in the word pointed to by the value of the left hand operand. As a result the effect of the Algol assignment statement "A:=B+C" is identical to the Bliss assignment "A\infty.B+C". This interpretation of names also allows the computation of pointers in Bliss so that the effect of the assignment "(A+3)\infty.(A+5)" is to store the value of the fifth location past A into the third location past A.

CONTROL STRUCTURES

Bliss is a block-structured, go-to-less, "expression language". That is, every executable construct, including those which manifest control, is an expression and computes a value. Expressions may be concatenated with semicolons to form expression sequences. An expression sequence is evaluated in strictly left-to-right order and its value is that of its last (rightmost) component expression. A pair of symbols, begin and end, or left and right parentheses, may be used to embrace such an expression sequence to form a simple expression. A block is a special case of the construction which contains declarations.

Other than expressions and functions, control mechanisms in Bliss fall into four classes: conditional, selection, looping, and leave. The conditional expression

if eo then e1 else e2

is defined to have the value e_1 just in the case that e_0 evaluates to \underline{true} and e_2 otherwise. The abbreviated form "if e_0 then e_1 "is considered to be "if e_0 then e_1 else e_1 ".

The conditional expression provides two-way branching. The <u>case</u> and select expressions provide n-way branching:

The <u>case</u> expression is executed as follows: (1) the expression e_0 is evaluated, (2) the value of e_0 is used as an index to choose one of the e_j 's ($1 \le j \le n$). The value of e_0 is constrained to lie in the range $1 \le e_0 \le n$. The value of the <u>case</u> expression is e_i ($i = e_0$). The <u>select</u> expression is similar to the <u>case</u> expression except that e_0 is used in conjunction with the $e_{2,j-1}$'s to choose among the $e_{2,j}$'s. The execution of the <u>select</u> expression above is described by the following, equivalent Bliss expression.

 $(\mathsf{T} {\leftarrow} \mathsf{e}_0; \ \mathsf{V} {\leftarrow} {-} \mathsf{1}; \ \underline{\mathsf{if}} \ \mathsf{e}_1 \ \underline{\mathsf{eql}} \ . \mathsf{T} \ \underline{\mathsf{then}} \ \mathsf{V} {\leftarrow} \mathsf{e}_2; \ ...$

if e2n-1 eal .T then V←e2n; .V)

Hence the value of the <u>select</u> expression is that of the last e_{2j} to be executed or -1 if none of them is executed.

The loop expressions imply repeated execution (possible zero times) of an expression until a specific condition is satisfied. There are several forms, some of which are:

do eo while e1

incr <id> from e0 to e1 by e2 do e3

In the first form the expression e_0 is repeated so long as e_1 satisfies the Bliss definition of true. The second form is similar to the "step ... until" construct of Algol, except (1) the control variable, <id>, <id>, is local to the incr expression, and e_0 , e_1 , and e_2 are evaluated only once (before the evaluation of the loop body, e_3). Except for the possibility of a leave expression within e_3 (see below) the value of a loop expression is uniformly taken to be -1.

The control mechanisms described above are either similar to, or slight generalizations of constructs in many other languages. Of themselves they do not remove the inconveniences generated by removing the goto. Another mechanism is desirable -- the leave mechanism. A leave is a highly structured form of forward branch which is constrained to terminate coincidentally with some control environment in which the leave is nested. The general form is:

leave <label> with <expression>

where <label> must be attached to a control environment within which the leave expression is nested. A leave expression causes control to immediately exit from a specified control environment. The expression> defines the value of the environment.

Finally, functions are defined and called in Bliss in a manner similar to that in Algol, except that there are no specifications and all parameters are implicitly call-by-value. The value of a function is the value of the expression forming its body.

CHAPTER II

OPTIMIZATION PRIMITIVES

This chapter develops a set of primitive relations, functions, and operators to be used in defining a class of feasible object code optimizations. There are several goals that direct this development.

First, the primitives are to form a basis for a set of <u>concise</u> descriptions of various optimizations. The compact notation of the system of primitives provides a basis for succinct descriptions of optimization strategies which in the past have often been described by lengthy algorithms.

Second, the primitives make possible a uniform representation of a large class of optimizations. The pyramid effect resulting from a buildup of primitives defined in terms of combinations of more basic primitives creates this uniformity. In addition this buildup produces a common basis for describing a wide range of optimizations.

Finally, the collection of primitives must allow an implementation of optimizations which is as concise as their descriptions. This final goal directs the selection among a number of different sets of primitives satisfying the preceding two criteria.

PRIMITIVE ORDERING RELATIONS

The problem of object code optimization can be viewed as the search for a transformation T which when applied to a program P produces an program P' that is more efficient. In general the optimization of a program effects a trade-off among a number of measures of program "efficiency". The most important include: size of the object code, execution time, and the amount of storage for data including user requested space and compiler generated temporary storage. The primitives presented this thesis will concentrate on exposing the set of feasible optimizations in a program. Even though a particular aspect of a program could be optimized (i.e. feasible), it may not be desirable because it only moderately decreases one of the above measures while increasing the It should also be pointed out that the notion of efficiency for an algorithm P cannot always be divorced from the data on which P executes. The optimization strategies to be considered and the primitives to be developed are in the class of data independent Data sensitive compile time. that are realizable optimizations optimizations in general require the collection of run-time statistics which can be used subsequently in re-compilation of the program. As the various optimization strategies are described their effect on the measures listed above will be noted.

We approach the problem of describing feasible optimizations for a program P by considering the ordering relations inherent in a

representation of P. There are several: the lexical order of the input text, the precedence-induced order of evaluation, both data-sensitive and data-insensitive order induced by control flow, a leftmost, depth-first tree order, and so forth. Two such orderings are of interest to the development.

The first is the order relation that results from considering a program as a mapping from its set of input variables to its set of output variables. Stated another way, this ordering, called the essential ordering and symbolized by "¬", is the ordering on evaluation of expressions that results from the application of the data flow and control. flow semantics of a language L to the set of expressions E in a program P. The optimizations to be considered will regard the essential order in a program as immutable.

The second ordering to be defined allows the selection of subsets of the total set of expressions in a program which at a given point are of interest to an optimization strategy. The following set of examples helps motivate the particular definition given for Bliss.

A representation of a program defines (at least partially) an evaluation order on its set of expressions. For example, the compound expression

 $\frac{begin\ e_1;\ e_2; \dots; e_n\ \underline{end}}{defines\ an\ ordering\ implying\ that\ evaluation\ of\ e_1\ precedes\ evaluation\ of\ e_2\ precedes\ evaluation\ of\ e_3\ precedes\ evaluation\ of\ e_4\ precedes\ evaluation\ of\ e_5\ precedes\ evaluation\ of\ e_6\ precedes\ evaluation\ of\ e_7\ precedes\ evaluation\ of\ e_8\ precedes\ evaluation\ of\ e_9\ precedes\ evaluation\ ev$

e₂ and so on. However the ordering inherent in this particular representation may or may not correspond to the ≺-ordering. The ≺-ordering might allow a number of permutations of the components of this compound expression. Consider the expression

e1 + e2.

While the <-ordering requires that the evaluation of e₁ and e₂ precede the evaluation of the sum, some languages may not define the <-ordering between the evaluation of the operands e₁ and e₂.

The initial ordering on a program is symbolized by "-d". Intuitively the relation e d e' expresses the notion that in a straightforward evaluation (i.e. that performed by a classical one-pass, non-optimizing compiler) of this representation of the program the evaluation of e would necessarily have preceded the evaluation of e'. This ordering reflects the precedence relationships of the program as exemplified in the addition expression above. It also reflects the sequential nature of execution as in the case of the compound expression. It does not, on the other hand, necessarily reflect the subnode relationship between nodes. Again, it is to be emphasized that the purpose of this ordering relation is to enable us to select subsets of expressions over which particular optimization strategies will operate.

Definition

The <u>initial ordering</u> on the set of expressions E of a Bliss program is defined as follows: Let e be a well-formed Bliss expression.

Define S(e)={e' ∈ E: e' ⊲ e and e' is a subexpression of e} U {e}.

One of the following cases applies for e:

- (1) e₁ <binop> e₂: e₁ ⊲ e, e₂ ⊲ e
- (2) <unop> e₁: e₁ < e
- (3) begin e1; ...; en end: e; ⊲ S(e;+1) (1≤i<n), en ⊲ e
- (4) case e₀ of set e₁; ...; e_n tes: e₀ ⊲ e, e₀ ⊲ S(e_i) (1≤i≤n)
- (5) if e0 then e1 else e2: e0 4 S(e1), e0 4 S(e2), e0 4 e
- (6) <u>select</u> e_0 <u>of</u> nset $e_1:e_2; ...; e_{2n-1}:e_{2n}$ <u>tesn</u>: $e_0 \triangleleft e, e_{2i-1} \triangleleft e \ (1 \le i \le n), \ e_0 \triangleleft S(e_{2i-1}) \ (1 \le i \le n),$ $e_{2i-1} \triangleleft S(e_{2n}) \ (1 \le i \le n)$
- (7) while e₁ do e₂: e₁ ላ S(e₂)
- (8) do e₁ while e2: e1 4 S(e2)
- (9) incr | from e1 to e2 by e3 do e4:
 - e₁ d e₂ d e₃ d e, e₁ d S(e₂), e₂ d S(e₃), e₃ d S(e₄)
- (10) $e_0(e_1, ..., e_n)$: $e_i \triangleleft S(e_{i+1})$ (0 $\leq i < n$), $e_n \triangleleft e$
- (11) <u>leave</u> <label> with e1: e1 4 e.

Consider the following piece of program text:

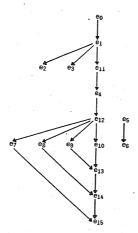
e₀; <u>if</u> e₁ <u>then</u> e₂ <u>else</u> e₃; e₄; <u>do</u> e₅ <u>while</u> e₆; e₇←e₈+e₉*e₁₀

where e₀, ... ,e₁₀ ∈ E. In addition define:

e11: if e1 then e2 else e3, e12: do e5 while e6;

e13: e9 * e10, e14: e8 + e13, e15: e7 ← e14.

The following lattice illustrates the total set of 4-relations that hold among the expressions e_0 , ... , e_{15} . ($e_i \triangleleft e_j$ if there is a path downward from e_i to e_j).



As the set of primitives continues to emerge, we will point out more detailed motivation for some components of the 4-ordering definition.

DECOMPOSITION OF ≺-ORDERING

In the case of simple non-control expressions such as e_{15} the \triangleleft -ordering reflects the precedence-induced ordering of the binary operation. For example the expression e_{13} initially precedes e_{14} . The

same relation, i.e. $e_{13} \prec e_{14}$, held in the essential ordering. The differences between the initial and essential orderings must be examined in more detail.

Most languages contain control environments whose components are potentially 4-order independent. Consider the following compound expression:

where A, ... ,F are distinct variables. Certainly the ⊲-order cf execution of these three assignments can be altered. For example

produces the same effect. Even within the context of a simple expression such as

the commutativity of the "+" operator is reflected in the fact that the development of the two multiplications is not defined. Nevertheless, the defore still reflects the requirement that both products be evaluated before the addition. Considering a third case.

$$.A * .B + F()$$

is an expression where the semantics of Bliss and Algol differ. In the case of Bliss neither the 4-order nor the 4-order of evaluation of the product ".A * .B" and the function call "F()" is well-defined. The usual interpretation of the semantics of Algol, on the other hand, imposes a strict left-to-right evaluation in the presence of the potential

side-effects resulting from the call on F. Our observations to this point on the 4-order and 4-order can be summarized by noting that in general the 4-order is weaker than the 4-order, i.e. e 4 e' implies that e 4 e' whereas the converse does not necessarily hold. That is, in some instances, the fact that e has been placed "before" e' (in the 4 sense) by the programmer is essential and sometimes it is not.

The optimization strategies discussed below will alter the 4-ordering in a program. Since the validity of such an alteration is constrained by the 4-ordering, a means must be provided for expressing the validity of transformations of a program. Given a pair of expressions e, e' where e 4 e', two aspects of the essential ordering can be identified that decide the validity of an optimizing transformation re-ordering e and e'. The first of these orderings reflects the relationship between an expression e and those of its subexpressions essential to its evaluation.

Definition

Let e_1 , $e_2 \in E$. e_1 is a <u>necessary constituent</u> of e_2 (notation: e_1 < e_2) if and only if (iff)

(1) e1 is a subexpression of e2, and

(2) evaluation of e2 requires prior evaluation of e1.

At first sight conditions (1) and (2) above may appear redundant. Indeed, if the language is Algol, they are redundant. However, in an expression language like Bliss, the following example illustrates their non-redundancy.

Example

Let e_1 : .A*.B, e_2 : e_1 +.C, e_3 : $D \leftarrow e_1$, e_4 : $(e_3; e_2)$. Then the following relations hold: e_1 < e_2 , e_1 < e_3 , e_2 < e_4 , e_3 \$ e_4 .†

Notice that the <-relation reflects a relationship only between values of expressions. In the example above the existence of e4 in a program requires that e3 be executed at some point. However e3 £ e4 indicates that the value of e4 can be computed without prior computation of the value of e3. The second ordering related to the essential ordering deals with the issue of side effects.

Definition

Let $e_1,\ e_2\in E.$ The expression e_1 is an essential predecessor of e_2 (notation: $e_1\ll e_2$) iff

(1) e₁ ⊲ e₂

(2) the evaluation of the sequences {e1,e2} and {e2} ({e2,e1} and {e1}) produce distinct values for e2 (e1).

Example

Let e_1 : A+A+1, e_2 : C+B*A+D, e_3 : E+A+(A+A+1), e_4 : D+B*C, e_5 : B*A. The following relations hold in the context of the compound expression: $(e_3;\ e_2;\ e_4)$. $e_1 \ll e_2$, $e_1 < e_3$ and $e_1 \ll e_3$, $e_1 \not\leftarrow e_4$, $e_5 < e_2$ and $e_5 \not\leftarrow e_2$.

 ⁺ Uniformly, a slash through a relational operator denotes the complementary relation.

their standard mathematical representations as subsets of ExE, then their relationship can be stated as: $\{\prec\} \subset \{<\} \cup \{\ll\}$. Hence it follows that if e < e' or e << e' then e < e', or equivalently if e > does not precede e' in the \prec -ordering then $e \neq e'$ and $e \neq e'$.

This section concludes by defining a relation on ExE which makes some of the subsequent discussions more convenient. <u>Independent</u> expressions are those whose <-ordering is not determined by the semantics of the language.

Definition

Let $e_1,e_2 \in E$. e_1 is independent of e_2 (notation: $e_1 \Leftrightarrow e_2$) iff $e_1 \not\in e_2, e_2 \not\in e_1, e_1 \not\in e_2, e_2 \not\in e_1$.

The usefulness of these primitive relations will become apparent during the discussion of the classical optimization strategies involving code motions.

SIMILARITY FUNCTIONS -- AN INTRODUCTION

Another primitive notion to be used in defining optimization strategies is a class of real-valued functions defined on the domain ExEcalled <u>similarity</u> functions. Chapter III will contain a more detailed discussion.

First, we introduce an equivalence relation called <u>congruence</u> on ExE which is an extension of the equality relation on E. Intuitively, two

expressions are congruent if there exists a one-to-one correspondence between them that preserves the tree structure and in which the corresponding nodes are identical operators or terminals. More precisely, the elements of E, considered as nodes in the tree representation, can be decomposed into non-terminal (N) and terminal nodes (T). Moreover T itself can be decomposed into names and literals. Recalling from our description of the tree representation of an expression in Chapter I that a node in E specifies its operator and operands, the notion of congruence is defined as follows.

Definition

Let e, e' \in E. e is congruent to e' (notation; e \cong e') iff either

(1) e, e' ∈ N,

e[operator] = e'[operator],

e[# of operands] = e'[# of operands]=n, and e[operand;] \cong e'[operand;] (1 \le i \le n);

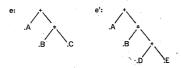
(2) e, e' ∈ T,

e and e' are equal literals or identical names.†

The idea for the similarity function grows out of the recognition that common subexpressions, which among other characteristics are congruent, are a rich source of optimizable expressions in a program. This observation suggests the consideration of those expressions which are "almost" common subexpressions but fail only because they are not quite congruent. As an

[†] The statement that e and e' are identical names is stronger than character string equality. Here we mean that they in fact refer to the unique variables accessible by that identifier within the present environment.

example, let e: .A+B*.C and e': .A+B*.(D+.E). When viewed in the form of the tree representations of e and e':



a strong correspondance is noticeable in the overall super-structure of the expression trees. The intuitive notion, then, of a similarity function is that it is a measure of how "close" two expressions come to being congruent. This intuition leads to the following minimal requirements for a function to be considered a similarity function:

σ is a similarity function only if

- (1) $\sigma: ExE \rightarrow [\emptyset, \infty)$,
- (2) $\sigma(e_1,e_2)=\emptyset$ iff $e_1\cong e_2$, and
- (3) $\sigma(e_1,e_2) = \sigma(e_2,e_1)$ for all $e_1,e_2 \in E$. (symmetric)

These three requirements should elicit the intuition that σ satisfies the requirements of a metric on E. That is indeed a reasonable intuition. What is not clear at this point is whether the additional metric requirement of the triangular inequality would add anything to the notion of a similarity function. It is clear, on the other hand, that the above restrictions are not sufficient to provide in themselves a very interesting class of functions. Further discussion of the characteristics of this

class of functions is deferred to Chapter IIi.

Given a particular similarity function, σ , a parameter δ (very σ -dependent) can be selected in terms of which the following relation on EXF can be defined

Definition

e is strongly similar to e' (notation: e \approx e') iff σ (e,e') < 8. †

The primary reason for interjecting this brief introduction to the similarity function at this point has been to establish the interconnection between the concepts of congruence, strong similarity, and the notion of similarity function. T roughout the remainder of this chapter the similarity function will be used in the very restricted sense of congruence. Chapter III will concentrate on exposing the overall motivation and usefulness of the concept.

This section concludes by introducing one more relation on ExE. Later on in Chapter II there will be a discussion demonstrating how the set of "redundant computations" as defined in Cocke and Schwartz are exposed by our primitives. For the present the notion of common subexpression, which specifies a subset of the collection of all redundant expressions, is defined in terms of the primitives developed so far.

[†] Local context will be sufficient to distinguish the use of "<" as a symbol for "less than" and for "necessary constituent".

Definition

e and e' are common subexpressions (notation: e = e') iff

 $(1) e \cong e',$

(2) e ⊲ e' or e' ⊲ e, and

(3) assuming e ⊲ e', ∀ e" such that e ⊲ e" ⊲ e', e ≮ e".

The intuition to be conveyed by this definition of a common subexpression is that if e = e', then (1) the values returned from the evaluation of e and e' are always identical and (2) the control flow of P is such that whenever e' is evaluated then e has been evaluated prior to it (or vice versa). The components of the definition mirror this intuition by saying that (1) e and e' are congruent, (2) the evaluation of e initially precedes e' (or vice versa) by definition of the 4-ordering, and (3) all the expressions that intervene between e and e' have the property that they do not produce side-effects that affect the value of e (equivalently: e') nor does e produce side effects on them. The latter condition says intuitively that the evaluation of e' is unnecessary since its value is available from the evaluation of e.

CODE MOTION OPTIMIZATIONS

The literature on object code optimization in the presence of control flow identifies a collection of optimization strategies called code motions. This set of optimizations falls into two subcategories: (1) moving evaluations of expressions to less frequently executed points in the

program and (2) avoiding unnecessary re-evaluation of expressions whose component values have not changed. The definition of common subexpression in the preceding section is an example of category 2.

CODE MOTION IN LINEAR BLOCKS

The collection of code motion optimizations about to be described are all predicated on a recursive, inside-out approach for their detection. For example, in detecting code motions relative to an if-then-else control environment, the detection proceeds by first invoking the optimization on the "then" and "else" expressions. The optimization on each of these expressions will (1) detect the feasible optimizations within its own local environment, and (2) return information to be used in detecting optimizations relative to the if-then-else environment. This overall approach requires that a means be provided for stating precisely what information about the sub-components of a control expression is required in order to detect optimizations for the control expression itself. The notion of a linear block is introduced for this purpose. Roughly speaking, β corresponds to those subexpressions of e through which a linear (i.e.

Definition

Let $e \in E$ and $E'=\{e' \in E: e' \text{ is a subexpression of } e\}$. The $\underline{\text{linear}}$ $\underline{\text{block}}$ β relative to e (notation: $\beta|e|$ is the set $\beta|e|=\{e' \in E': e' \prec e\}$.

Since in the context of the use of $\beta|e$ the expression e is quite often obvious, "|e" is simply omitted in most cases. By convention, the linear block relative to e_i will be denoted by β_i . In flow diagrams, linear blocks are depicted as unbroken vertical lines (flow passing from top to bottom):

β

Example

Consider the expression:

 $\begin{array}{l} & \text{if } e_0 \\ & \text{then } (e_1; \ do \ e_2 \ \text{while } e_3; \ e_4) \\ & \underline{e|se} \ (e_5; \ \underline{if} \ e_6 \ \underline{then} \ e_7 \ \underline{e|se} \ e_8) \\ \text{and define:} \\ & e_9: \ \underline{do} \ e_2 \ \underline{while} \ e_3, \\ & e_{11}: \ (e_1; e_9; e_4), \\ & \text{and } e_{12}: \ (e_5; e_{10}). \\ \end{array}$

Consider a linear block β that contains the element e: A+B*C. We wish to develop a concise description of the potential movability of e backward (to the top) or forward (to the bottom) of the block. It may be feasible to move the evaluation of B*C backwards even though the entire expression e cannot be moved. For example:

(F(.A): A←.B*.C...

Assuming F does not produce side effects on B or C, we recognize that the evaluation of .B*.C can be moved backward to the head of the linear block

whereas the store into A must follow the parameter evaluation for the call on F. In our terminology, the expression F(A) is an essential predecessor of e but not of .B*C. On the other hand the evaluation of .B*C can never be moved forward to a point after the evaluation of e since .B*C is a necessary constituent of e.

The following definition defines three sets which make the succeeding definition less cumbersome.

Definition

```
Let e \in \beta, \beta a linear block. pro-dominator (\beta,e) = \{e' \in \beta: e' \triangleleft e, e' \ll e \text{ or } e \ll e'\}, epi-dominator (\beta,e) = \{e' \in \beta: e \triangleleft e', e \ll e' \text{ or } e' \ll e\}, post-dominator (\beta,e) = \{e' \in \beta: e \triangleleft e', e' \not \Rightarrow e\}.
```

The pro-dominator set contains those elements of β which initially precede e such that they produce a side effect on e or e produces a side-effect on The epi-dominator set differs from the pro-dominator only in that pro-dominator initially follow Intuitively the its elements (epi-dominator) contains those elements of β which prevent the movement of e backward (forward) to the head (tail) of β because they produce a side-effect on e or vice versa. The post-dominator set consists of those elements of β which initially follow e and are not independent of e. Hence the post-dominator consists of those elements which prevent the movement of e forward either because of a side-effects relationship or because their evaluation requires the evaluation of e. It follows from the definitions of " \ll " and " \diamond " that: epi-dominator(β ,e) \subseteq post-dominator(β ,e).

<u>Definition</u>

Let β be a linear block. prolog(β) = $\{e \in \beta: \text{pro-dominator}(\beta,e) = \emptyset\}$, epilog(β) = $\{e \in \beta: \text{epi-dominator}(\beta,e) = \emptyset\}$, postlog(β) = $\{e \in \beta: \text{post-dominator}(\beta,e) = \emptyset\}$.

Note that it follows immediately that $postlog(\beta) \subseteq epilog(\beta)$.

Example

Let e: (Ac.B; if A gtr B then Cc.B*C; Dc.C; Bc.X*Y; X-3). Then $\beta = \{B, Ac.B, A, B, A gtr B, C, Dc.C, X, Y, X*Y, Bc.X*XY X-3\}.$ Note that in our discussions of code motion and the related subsets of linear blocks, constants (names and literals) will not be listed since they do not enter into the feasibility of code motions. Now:

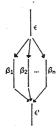
Observe that the second .B in $prolog(\beta)$ is the right operand of .A gtr .B and that the .C in $epilog(\beta)$ is the right side of the store .De.C.

These sets define those expressions that can be moved forward or backward relative to the head or tail of β. At this point the utility of these sets is not yet clear but their usefulness becomes apparent in the context of control environments. In particular the next two sections on optimization strategies for branching and loop control environments stress the expressive power of the primitives for generating concise descriptions of a variety of optimizations.

OPTIMIZATION PRIMITIVES CODE MOTION IN LINEAR BLOCKS

CODE MOTION IN FORKED CONTROL

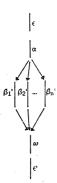
Consider a branching control construct of the form



where € functions as a selector among the n branches. This form of control represents both if-then-else and case types of control environments. The following sections describe several optimization strategies relative to this form of control environment.

ALPHA-OMEGA CODE MOTIONS

The first form of feasible optimization exploits the code motions of the linear blocks β_1, \dots, β_n so that the following flow diagram results:



The linear blocks α and ω contain those expressions factored forward and backward from all of the branches, β_1 .

Example

Let β_1 : $(A \leftarrow X \approx Y, Y \leftarrow 3)$ and β_2 : $(B \leftarrow X \approx Y, Y \leftarrow 3)$. Consider the expression: if ϵ then β_1 else β_2 . A feasible optimization is to let α : $T \leftarrow X \approx Y, \alpha$: $Y \leftarrow 3$, β_1 : $A \leftarrow T$, β_2 : $B \leftarrow T$. This yields the expression: (if $(T \leftarrow X \approx Y, \epsilon)$) then $A \leftarrow T$ else $B \leftarrow T$; $Y \leftarrow 3$).

A primary goal of the development of our optimization primitives is to provide a means of concisely describing the set of feasible members of sets such as α and ω . To that end an operator on the power set of E is introduced.

Definition

While formal intersection is different from ordinary set intersection the analogy should be obvious: formal intersection differs from set intersection in that the equivalence relation of equality of elements is replaced by that of congruence.

Example

Let β_1 and β_2 be as defined in the preceding example. Then $\beta_1 \wedge \beta_2 = \{X, Y, X * Y, Y * 3\}$. We reinforce the fact that the "\n" operator differs from ordinary set intersection by noting that $\beta_1 \cap \beta_2 = \emptyset$.

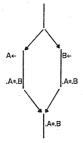
The notion of formal intersection provides us with a powerful tool to concisely define the sets α and ω .

Given a forked control environment with branches e_1,\dots,e_n , the domain of elements (α) available for pre-evaluation is described by: $\alpha\subseteq \wedge$ prolog (β_1) . The domain of elements (ω) available for post-evaluation is described by: $\omega\subseteq \wedge$ postlog (β_1) .

The optimizations described by the sets α and ω are examples of optimizations that save space, do not effect time, but may prolong the life-time of temporary storage locations.

POST-MERGE RE-EVALUATIONS

In addition to the goal of providing a collection of primitives that allow a concise definition of a variety of optimizations, these primitives should also be "complete" in the sense that they may be used to describe the class of "redundant" computations in a program. Consider the following example:

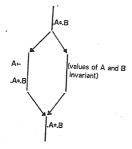


It is apparent that the product that follows the merge point need not be re-evaluated. The set of expressions available for this optimization is described by:

$\land epilog(\beta_i).$

Recall that an element of the epilog set cannot in general have its evaluation moved to the end of the linear block. The value, however, of such expressions is not altered by the expressions that succeed them in the block.

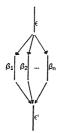
In practice, a more general case can be considered. For example:



once again the evaluation of the product after the merge is not necessary.

Since .A*.B does not appear in the right hand branch, it would not be an element of the formal intersect of the epilog of the branches. The extension is straightforward.

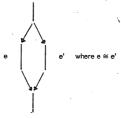
Consider



Given a forked control environment with selector expression ϵ and branches e_1, \dots, e_n , define $e_i' = (\epsilon; e_i)$ and $\beta_i' = \beta_i|e_i', 1 \le i \le n$. Then the set of expressions whose evaluation at the merge point would be redundant is the set: $\land epilog(\beta_i')$.

WASP-WAISTING

There is another class of optimizations one might consider in a branching control environment. Consider the example:



Assume the unspecified portions of the branches are such that the

simultaneous motion of e and e' backward as well as forward is impossible.

One can consider an optimization which because of the altered appearance of the flow diagram, we call "wasp-waisting".



The jump to and return from the common evaluation point can be accomplished by subroutine call and return instructions or by re-testing the branch condition. Wasp-waisting turns out to be a particular case of a more general class of optimizations using the notion of similarity, which will be discussed in Chapter III.

CODE MOTION IN LOOPS

The looping constructs to be considered consist of a body β_1 and a predicate β_2 to be evaluated on each iteration. This section will consider two types. A "do-while" form has its test at the bottom of the loop.



A "while-do" form has its test at the top of the loop.



Other forms of loops such as counting types can be modeled by these forms.

LOOP INVARIANT EXPRESSIONS

The first optimization strategy considered is the pre-evaluation of the "loop invariant expressions", i.e. those whose values do not change on any iteration of the loop. In terms of the primitives developed, the description of the set of loop invariant expressions is straightforward.

Given a loop control environment, the set of loop invariant expressions is described by: prolog(β) Ω epilog(β), where β is the linear block relative to the compound expression (β ₁; β ₂) in the "do-while" case and (β ₂; β ₁) in the "while-do" case.

The description has intuitive appeal since it simply states that any

expression whose evaluation is not affected by occurring either before or after the loop is not changed by execution of the loop.

CYCLIC RE-EVALUATIONS

The cyclic nature of loop control gives rise to a particular class of "redundant" computations. Consider the following example:

.A*.B A--.A*.B

Clearly the expression .A*.B is not invariant throughout the loop. However if the expression .A*.B were pre-evaluated at entry to the loop and stored in a temporary T and if after each recomputation of A or B the expression .A*.B were again evaluated in T, there would be no need to re-evaluate .A*.B at the top of the loop on each iteration. The restructured computation is:



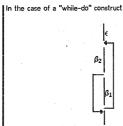
Given a loop control environment where β is the linear block relative to the expression (β_1,β_2) ("do-while") or (β_2,β_1) ("while-do"), the set of expressions whose evaluations at the head of β are redundant to avaluations at the tail of β are described by the set: $\operatorname{prolog}(\beta)$ δ -pilog (β) .

Comparison of this set with the set of loop invariant expressions described above reinforces the distinction between the notions of formal intersection and set intersection. In the case of a loop invariant expression e, the expression e itself appeared in both the prolog and epilog sets whereas an element of the formal intersect is simply an expression which has a formally identical image in both the prolog and epilog sets. Since the first instance of A&B can be moved backward, it appears in the prolog but the redefinition of A prevents its appearance in the epilog. However the second instance of A&B can be moved forward and

so appears in the epilog.

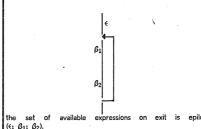
POST LOOP RE-EVALUATIONS

Finally, let us point out how loops participate in the exposure of the set of redundant expressions to their surrounding environment.



the set of expressions whose values are available on exit from the loop is the set epilog(ϵ ; β_2).

For the case of a "do-while" construct



CODE MOTIONS AND LEAVE EXPRESSIONS

To this point our discussion has been limited to go-to-less control structures. In this section we consider the effect of introducing the Bliss "leave" mechanism for exiting control environments. In particular, are the set of primitives powerful enough to describe the code motion optimizations in the presence of leave expressions?

Consider the following example:

eo;

LOOP: while e1 do

(e2; if e3 then leave LOOP with e4; e5; e6); e7

The flow diagram for this expression is:

The class of code motion optimizations we have been discussing can be divided into three subclasses: (1) moving an evaluation backward, (2) moving an evaluation forward, or (3) eliminating an evaluation because it is available on all control paths leading to the present evaluation. The linear block relative to the leave expression participates in optimizations of class (1) in a manner analogous to the optimizations proposed for an expression of the form:

if e3 then e4 else (e5; e6).

As for optimizations of the classes (2) and (3), it is analogous to optimizations for the expression:

if <arbitrary predicate> then e4

else while e1 do (e2; e3; e5; e6).

In effect a leave expression is a forking construct whose optimizations involving backward motion of code behave as though the fork is local to the environment surrounding the leave and whose optimizations involving forward motion of code behave as though the fork terminates at the termination point of the labelled expression. As a particular example, the set of expressions whose evaluation is available on termination of the LOOP expression above is the set: epilog(eq:e4) A epilog(eq:e1).

STRENGTH REDUCTION

A classical optimization in the presence of iterative loop control is "strength reduction". Basically strength reduction exploits the inductive behavior of the control variable in a loop in the attempt to replace relatively expensive operations with less expensive ones by applying recursion relations to express the expensive operation in terms of the less expensive one.

The following example illustrates the technique. Assume a segment of storage named A has been structured so that access to the I-th element of A is defined by the Bliss expression: A+3*I+5. The loop which follows will zero out every 3*K-th element of A starting at the (3*M+5)-th element and ending at the (3*N+5)-th element of A.

incr | from .M to .N by .K do (A+3+,I+5)←0

that on each iteration of the loop the relation

Note that on each iteration of the loop the relatively expensive multiplication 3*. I must be re-evaluated in the loop body.

Strength reduction on such a loop transforms the loop expression above to the following:

incr | from (A+3*.M+5) to (A+3*.N+5) by 3*.K do .I+0.

This latter loop has the same semantic effect as the former but now there are no multiplications taking place in the loop body.

Unrolling the first few iterations of the original loop will help motivate the discussion which follows.

- (Ø) I←,M
- (1) if .l gtr .N then <endloop>;
- (2) (A+3*.I+5)←Ø;
- (3) I←,I+,K;
- (4) if .1 gtr .N then <endloop>;
- (5) (A+3*.I+5)←Ø;
- (6) I←.I+.K;

etc.

Notice that the accessing expressions in (2) and (5) are congruent. They are not common subexpressions, however, because of the intervening re-evaluation of I at (3). This unrolled representation of the loop example suggests an investigation into a more general form of the strength reduction notion.

STRENGTH REDUCTION -- A GENERALIZATION

Consider the following question: given e, e' \in E and e \triangleleft e', can one characterize the cases in which there exist an expression Δ e and a function F such that e' \equiv F(e, Δ e) and the computational cost of F(e, Δ e) is less than the computational cost of e' evaluated in the usual manner? We have already seen two cases:

(1) Clearly the example of a strength reduction optimization in the preceding section fits this situation. In general it reduced

the cost of execution.

(2) The second case involves the redundant expression elimination discussed earlier in the chapter. The sequence (X \leftarrow e; Y \leftarrow e <op> e') will make use of the fact that it need not recompute the left hand operand of <op>. Such optimizations save both time and space.

Our discussion of strength reduction examines the possible extensions of the notion and the corresponding difficulties in exploiting those extensions. In the process of this development the primitives already developed are used and a few specialized notions are defined.

PRIMITIVES FOR STRENGTH REDUCTION

Returning to the context of the unrolled strength reduction example presented above, the necessity of stating more precisely the interaction between the re-evaluation of I and the corresponding change in the value of A+3*,I+5 is evident. We begin by proposing a definition that describes the set of expressions involved in the evaluation of e, e' and all the expressions between them in the <-ordering.

Definition

Let e, e' \in E, e < e' and E'= {e" \in E: e < e" < e'}. The <u>interval</u> from e to e' (notation: int(e,e')) is defined as the set E' U {e \in E: e" \in E', e a subexpression of e"}.

Example

Let e:(e₁; e₂; e₃; e₄) where e₃: \underline{do} e₅ while e₆. Then int(e₂,e₄) = {e₂, e₃, e₅, e₆, e₄}. Similarly int(e₁,e₄) = with the linear block β [e, [e₁, e₂, e₃, e₄, e₅, e₆]. Contrast this with the linear block β [e, β [e] = {e₁, e₂, e₃, e₄} does not include e₅ and e₆ because of the definition of the d-ordering.

The notion of linear block is defined relative to a single expression. As a result it is impossible to talk about the linear block relative to an interval. This difficulty is resolved by defining the minimal expression containing an interval, which will be called the <u>cover</u> of the interval. In some cases the cover will itself be an expression in the program. Consider the case, however, where the expressions e and e' appear in a compound expression (as did the two instances of A+3+1+5 in the unrolled loop example above). For example let e":(e₁; e₁ e₂; e₃; e'; e₄; e₅) be the minimal expression containing e and e'. The interval from e to e' is the set {e, e₂, e₃, e'} and the cover should not contain expressions which will not enter into the consideration of what occurs as execution passes from e to e'. Hence the cover, in this case, will be defined as the compound expression (e; e₂; e₃; e') which does not occur as an expression in the program.

<u>Definition</u>

Let e, e' \in E, e $_{\circ}$ e', and let e" be the minimal expression in E which contains the elements of int(e,e') as subexpressions. The <u>cover</u> of int(e,e') is defined as e'', e'' not a compound expression

c, otherwise.
c is defined as follows. Let e":(e₁; ...; e₁; ...; e_j; ...; e_n).
Then c is the compound expression (e₁; ...; e_j) where:

 $\begin{array}{l} \forall \ x \in \text{int}(e,e') \ \exists \ k, \ i \le k \le j \ \text{such that} \ x \ \text{is a subexpression of} \ e_k, \\ \text{and} \\ \forall \ e_k, \ i \le k \le j, \ \exists \ x \in \text{int}(e,e') \ \ \text{such that} \ x \ \text{is a subexpression of} \ e_k. \end{array}$

Referring to the preceding example, it follows from the definition of a cover that $cover(e_1,e_3) = (e_1;e_2;e_3)$ and $cover(e_5,e_6) = e_3$.

The next concept is well understood but is defined for completeness.

<u>Definition</u>

Let $e_0, \dots, e_n \in E$ and let l_1, \dots, l_n be variables. A $\underbrace{\text{linear}}_{\text{polynomial}}$ e in the n variables l_1, \dots, l_n is denoted by $e < l_1, \dots, l_n >$ and is an expression of the form $e_0 + e_1 * l_1 * \dots * e_n * l_n = 0$

STRENGTH REDUCTION WITHOUT LOOPS

Now the conditions that make a strength reduction optimization possible in an environment such as the specific unrolled example above can be described. Let e, e' \in E, e \triangleleft e', and e \cong e'. Let e (and e') be linear polynomials in n variables: e<|1, ... , |_n> (e'<|1, ... , |_n>). Let β be the linear block relative to cover(e,e'). Assume that for all k, $1 \le k \le n$, the only redefinitions of 1_j (if any) in int(e,e') are of the form: $1_j \leftarrow 1_j + \Delta_j$ where $\Delta_j \in \text{prolog}(\beta) \cap \text{epilog}(\beta)$. Also assume that the coefficient expressions e_0, \ldots, e_n are elements of $\text{prolog}(\beta) \cap \text{epilog}(\beta)$. Define: $\Delta e = e < 1_1 + \Delta_1, \ldots, 1_n + \Delta_n > - e < 1_1, \ldots, 1_n >$. The following two observations can be made:

(1) Δe is a polynomial $(\Delta e < \Delta_1, ..., \Delta_n >)$ and moreover if $e = e_0 + e_1 * |_1 + ... + e_n * |_n$ then $\Delta e = e_1 * \Delta_1 + ... + e_n * \Delta_n$.

(2) value(e')=value(e)+value(Δe).

Example

Admittedly, the above example is biased by the fact that both the polynomial coefficients and Δ_1 and Δ_2 are all constants. If, for example, the re-definition of I were: I+.1+.K, then $\Delta e = 3*.K+28$. The product 3*.K is more palatable if we consider a sequence such as:

Now the evaluation of Δe occurs only once and the successive stores in the X_1 's can be accomplished by the sequence:

Δe←3*.K+28;

X₁←.e;

X₂←.X₁+∆e;

X₃←.X₂+∆e;

 $X_m \leftarrow X_{m-1} + \Delta e$

Assume that the sequence of names X_1 , X_2 , ... X_{m-1} , X_m is computable in the sense that a function f exists such that for all i, $2 \le i \le m$, $X_1 = f(X_{1-1})$. Then the sequence of stores above strongly suggests an unrolled loop.

STRENGTH REDUCTION WITH LOOPS

The observations made in the preceding section can be restated within the context of a looping control expression.

Given a loop of the form: incr I from e0 to e1 by e2 do e3, let e, a subexpression of e3, be a polynomial in I for which we wish to perform a strength reduction optimization. Let e<1> = e'*.1+e" and let β be the linear block relative to e3. Then the following conditions must hold for the strength reduction to be feasible:

- (1) e',e" ϵ prolog(β) Ω epilog(β), i.e. e',e" are loop invariants,
- (2) the only redefinition of I is the loop increment le-leg. (The semantics of Bliss require that eg's value be evaluated prior to loop entry and perserved. Hence eg is loop invariant.)

The strength reduction optimization is realized by transforming the original loop to the following:

 $\underline{\mathsf{incr}} \mid \underline{\mathsf{from}} \ (\mathsf{e_0}; \, \mathsf{l'\leftarrow}\mathsf{e}{<}\mathsf{e_0}{>}) \ \underline{\mathsf{to}} \ \mathsf{e_1} \ \underline{\mathsf{by}} \ (\mathsf{e_2}; \ \Delta \mathsf{e}{\leftarrow}\mathsf{e'*e_2})$

<u>do</u> (e₃'; l'←.l'+Δe);

where e_3 is obtained from e_3 by replacing e by .I'. If e were the only expression in e_3 that accessed the value of I then a more significant strength reduction of the form:

incr I from e<e0> to e<e1> by e'*e2 do e3'

can be performed where again e₃ is obtained by replacing e with I in e₃. This loop has only one induction variable and the "to" test on e₁ has been replaced by e<e₁>. The following section examines extensions of the strength reduction notion and the corresponding problems.

STRENGTH REDUCTION EXTENDED

incr | from Ø to N by 3 do F(3*.1*.1-4*.1+7);

to

```
i-7; T-3*.N*.N-4*.N+7; M-Ø;

<u>while</u> .l <u>leg</u> .T <u>do</u>

<u>begin</u>

F(.l);

|-.l+(.M+15);

M-.M+54;

end;.
```

In general if the expression e upon which a strength reduction is being performed is an n-th degree polynomial, then n-1 additional variables, like M in the example above, must be introduced in order to maintain the partial accumulations.

Having removed the linearity requirement for polynomials, consider the possibility of relaxing the polynomial requirement itself. The point of the reduction in strength optimization is to replace an expensive operation with a less expensive one. In the case of multiplication and addition, the feasibility of such a replacement comes from the inductive relationship between the operands of the successive multiplications and the fact that a product can be accumulated by a sequence of additions. This overall relationship is reflected in the fact that given a n-th degree polynomial e<1> then the polynomial e<1> e<1> is always of degree n-1. There are two critical points here:

- (1) Δe is a polynomial and so a closed form solution is available to the difference $e<1+\Delta_1>-e<1>$, and
- (2) Δe is of degree n-1, which means that successive reductions will eventually reduce all multiplications to additions.

Hence the question remains: are there other strength reductions besides those between "*" and "*"? All the preceding development holds equally well if we replace "*" by exponentiation and "*" by "*". For example the loop:

incr I from 1 to .N by 1 do A[.I]←.X <exp> .I;

can be replaced by the following expression in which no exponentiation occurs:

(J←.X; <u>incr</u> I <u>from</u> 1 <u>to</u> .N <u>by</u> 1 <u>do</u> (A[.i]←.J; J←.J*.J));

The section on strength reduction began by asking the question: given e, e' \in E and e \triangleleft e', can one characterize the case where e' \equiv F(e, \triangle e) and the cost of computing F is less than the cost of computing e'? The attempt to isolate the essential characteristics of strength reduction with a view to extending the notion initially motivated that question. discussion has pointed out two directions for extension: (1) strength reduction in non-looping environments and (2) strength reduction between The primitives were able to define non-polynomial expressions. feasible strength reductions independent of the presence of the loop control environment. The challenge remains in case (1) to discover an algorithm for directing the search through the set E for pairs (e,e') on which a strength reduction can be performed. Loops have the property that they both localize the search and, in the case of incr loops, immediately identify an induction variable. As for case (2), the challenge is to discover a more general means of constructing closed-form representations of the recursion relation F. Polynomials have the property that the expression Δe is an easily identifiable polynomial itself.

REDUNDANT EXPRESSIONS

Several references have been made in the preceding sections to the concept of redundant expressions in a program. In the present section we demonstrate that our primitives expose the set of redundant expressions in a program consisting of the forked and looping control environments discussed above. The following definition is a direct quote from the text by J. Cocke and J. Schwartz.†

Definition

An operation A=B (i.e. an operation which combines two inputs A and B to give some sort of result, which we write as A=B) is redundant if there exists no track in the program graph, either beginning at the program entry block, or beginning at any assignment of a new value to one of the variables A or B, which reaches the given operation without passing through some preceding calculation of the result A=B.

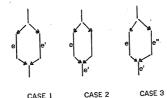
The definition of common subexpression identified a collection of redundant expressions, i.e. if e=e', then e' is redundant (assuming $e \triangleleft e'$). The fact that $e \triangleleft e'$ implies that every control path that leads to an evaluation of e' has previously evaluated e. There is no intervening assignment to the components of e' since by part(3) of the definition of common subexpression: $\forall e''$, e'', e'',

Assume that e' is a redundant expression and that e is the congruent expression that "creates" this redundancy. Furthermore, assume that this

[†] cf. [CS70], pp. 427-428.

redundancy was not exposed by the optimization techniques presented above. Now if e=e', then e' would be redundant. Hence one of the three conditions for a common subexpression must not hold. The first condition, viz. $e \cong e'$, must be satisfied by e and e' since congruence is a property of redundant expressions. If the second condition ($e \triangleleft e'$) is assumed to hold, then the third condition of the c-s-e definition indicates the existence of an expression e'' such that $e \not \sim e''$. However, the existance of the expression e'' again violates the definition of redundancy for e and e'. Thus we have only to consider the cases in which $e \not\sim e''$.

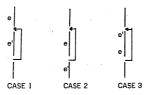
There are three cases to consider for forking control environments:



In case 1; the expression e' is not redundant since no control path leads from the left-hand branch to the right. Notice, however, that optimizations have been proposed above which attempt to combine the two evaluations. The α and ω optimizations expose the feasibility of simultaneously moving the evaluations of e and e' backward or forward. Wasp-waisting is a feasible optimization for those cases where forward or backward motion is impossible. In case 2, e' is not redundant since

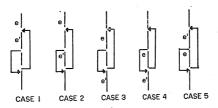
control potentially passes down the right-hand branch and so does not evaluate e. Finally in case 3, e' is potentially redundant since the expressions e and e" are evaluated on each control path. If no side-effect producing expressions occur between the evaluation points of e and e" and the evaluation point of e', then e' is redundant. This class of redundant expressions, as described in the section on post-merge re-evaluations, is detected by the formal intersection of the epilogs of the branches.

A "do-while" looping environment presents three cases to consider in which e & e':



For convenience, we define χ to be the set of loop-invariant expressions in a loop. Thus $\chi=\operatorname{prolog}(\beta)$ Ω epilog (β) where β is the linear block relative to the body and predicate of the loop. In case 1, the evaluation of e' is redundant only if $e' \in \chi$. The redundancy of e' in case 2 does not require that $e \in \chi$ but simply that $e \in \operatorname{epilog}(\beta)$. Case 3 is exactly the situation discussed in the section on cyclic re-evaluations. In this case e' is redundant if $e' \in \operatorname{prolog}(\beta) \land \operatorname{epilog}(\beta)$.

The "while-do" form of loop presents the following cases for consideration:



Again let χ be the set of loop-invariant expressions. Let β_1 and β_2 be the linear blocks relative to the <u>while</u> and <u>do</u> expressions respectively. Let $\beta = \beta|(\beta_1; \beta_2)$. In both cases 1 and 2, the expression e' is redundant only if $e' \in \chi$. In case 3, e' is redundant if $e' \in \text{epilog}(\beta_1)$. e' is not redundant in case 4 since there is no guarantee that the <u>do</u> expression will be executed. Finally, case 5 is again an example of a cyclic re-evaluation and so e' is redundant if $e' \in \text{prolog}(\beta) \land \text{epilog}(\beta')$.

SUMMARY

A primary goal of the thesis is to propose a collection of primitives for describing object code optimizations which are powerful enough to provide concise descriptions of optimizations. The set of primitives presented in this chapter was motivated, defined, and used in describing the code motion, redundant expression elimination, and strength reduction

optimizations discussed in Cocke and Schwartz. The collection of paragraphs delineated by vertical lines describe these optimizations. Their concision is self-evident.

The primitives also apply to a broad class of optimizations. In particular, it would be inappropriate that disjoint collections of primitives would be used in describing each class of optimizations. An examination of the set of descriptions shows that most of the primitives permeate through all the descriptions. The ordering relations $(4, \prec, \ll, \prec)$ and the subsets defined in terms of them (prolog, epilog, postlog) are used consistently throughout the chapter. As a result, although the optimizations themselves may on the surface appear to be unrelated, the primitives provide a homogeneous description of them. This homogeniety, in turn, leads to a compact, cleanly structured implementation.

Another objective of the thesis is that the primitives be language independent. This objective has been achieved by isolating the language dependent relationships in the "necessary constituent" (<) and "essential predecessor" (<<) relations. The ability to isolate these language dependent relationships contributes significantly to the concision of the descriptions.

The primitives have been developed in a representation-independent manner. No inherent characteristics of the primitives are concerned with the data structure of the program's representation. Hence there is no

implied implementation strategy underlying the primitives. Again, this contributes to their concision and clarity. This aspect of the primitives allows relative freedom in implementation strategies. In addition it has resulted in a set of primitives that can be manipulated purely on a formal level. Potentially, this can lead to results whose discovery would be hopelessly obscured by any specific representation.

object investigations in the area of Finally. optimizations often describe optimizations in terms of lengthy algorithms which manipulate particular representations. Our primitives have succeeded into operators, relations, and the partitioning those algorithms characteristic functions of particular sets of expressions. able to describe optimizations in terms of the primitives without regard to the representation of the program or the particular implementation details of the primitives. A good example of the effect of the homogeniety, concision, and representation-independence is the discussion of the completeness of redundant expression elimination in the preceding section.

CHAPTER III SIMILARITY FUNCTIONS

In Chapter II a collection of primitives was developed to concisely describe previously known optimization techniques. This chapter examines a class of real-valued functions called similarity functions to be used in conjunction with the primitives of Chapter II in describing a set of new optimizations. These optimizations produce dramatic reductions in object code size in certain cases where the classical optimizations presented earlier have little effect. In particular an example presented in Chapter IV shows a 28 percent savings in a 1000-word program resulting from these techniques. This reduction is to be contrasted with the 6 percent savings that results when the same program is optimized using only the classical optimizations.

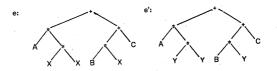
The presentation of similarity in this chapter is divided into three sections: (1) a discussion of the origins of the similarity concept, (2) the development of a particular similarity function, and (3) an examination of a collection of optimization techniques based on the concept.

ORIGINS OF THE SIMILARITY CONCEPT

The optimization techniques described in Chapter II fell into two categories: (1) moving the evaluation of expressions either to reduce frequency of execution (e.g. pre-evaluation of loop-invariant expressions) or to eliminate parallel evaluations (e.g. alpha-omega code motions) and (2) avoiding the unnecessary re-evaluation of an expression (e.g. a common subexpression). The initial stimulus for the similarity concept arises from a consideration of the sets of expressions on which optimizations from category (2) operate. If e and e' are a pair of expressions such that the evaluation of e' is made unnecessary by the prior evaluation of e, then e and e' are congruent and can be translated into identical code sequences. The phrase "identical code sequence" is to be interpreted loosely as meaning an identical sequence of machine code operations ignoring the possibility of different temporary accumulators. The key intuition is that although congruent expressions are translated into identical code sequences, the converse does not follow. That is, identical code sequences can be produced for the evaluation (or partial evaluation) of expressions which are not congruent. For example, the code sequence

can be used to evaluate e: A*X*X+B*X*C or e': A*Y*Y+B*Y*C by loading T_1 with X or Y respectively. In terms of the tree representations of

e and e':



the similarity of the code sequences produced for these two trees, and consequently the possibility of using a single sequence such as that above, arises from the common superstructure of the two trees. The notion of a similarity function is introduced precisely in order to measure the degree of identity of the superstructures of two trees. The similarity notion provides a coherent mechanism for identifying expressions whose evaluations can be merged into identical code sequences.

Additional intuition for the similarity concept is derived from a consideration of the requirements imposed on a pair of expressions e and e' by the definition of common subexpression. There are three: (1) e is congruent to e' (e≅e'); (2) e initially precedes e' (e¬de'); and (3) none of the expressions intervening between e and e' have a side-effects relationship with e (∀ e", e ⊴ e" ⊴ e', e ≮ e" and e" ≮ e).

The class of optimizations considered in Chapter II uniformly imposed condition (1). That set of optimizations was described in terms of formal intersection or ordinary set intersection. Because congruence is an

inherent characteristic of these two operators, those optimization strategies necessarily dealt with sets of expressions that were congruent. The same optimization techniques did, on the other hand, consider cases in which conditions (2) and (3) did not hold. In the collection of code motions related to forking environments, the sets α and ω contained expressions that did not satisfy condition (2) since in the initial-ordering of the program those expressions were on parallel branches and hence did not initially precede one another. The optimizations involving strength reduction and cyclic re-evaluations relaxed condition (3) by allowing the existence of intervening side-effects related expressions. Naturally enough, since several optimization strategies involved relaxation of conditions (2) and (3), one is led to consider relaxing condition (1).

As a framework for the ensuing discussion, we will present examples of optimization techniques involving a relaxation of condition. (1) which were not exposed by the primitives developed to this point. For example:

$$\begin{array}{c} \underline{if} \ e_0 \\ \underline{then} \ (e_1; \ ... \ ; \ e_k; \ A \leftarrow .B + .C * .D) \\ \underline{else} \ (e_{k+1}; \ ... \ ; \ e_n; \ A \leftarrow .B + .C * .E). \end{array}$$

Clearly both assignments to A can be evaluated by the code sequence

MULT T,C ADD T,B STORE T,A

where on the then and else branches T has been loaded with D and E respectively. Hence, an optimization strategy for this expression consists of replacing the expression with:

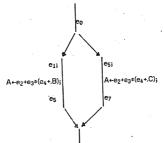
 $\begin{array}{c} (\underline{if}\ e_0\\ \underline{then}\ (e_1;\ ...\ ;\ e_k;\ T\leftarrow D)\\ \underline{e|se}\ (e_{k+1};\ ...\ ;\ e_n;\ T\leftarrow E);\ A\leftarrow B+.C*.T). \end{array}$

The ω set for forked control environments described in Chapter II did not expose this optimization since the assignment expressions are not congruent.

The preceding optimization technique depends on the fact that, although the expressions are not quite identical, they are very close to being identical. Hence one of the properties of a similarity function must In addition, the the quantification of this notion of "closeness". measurement must be sufficiently fine so that it can distinguish degrees of "closeness" rather than compute a simple boolean indicating "close" or "not These observations point out a major distinction between the similarity concept and the primitives developed in Chapter II. primitives served to expose the feasible optimizations in a program. the most part, the optimizations were not only feasible but also desirable since, in general, they reduced both object code size and execution time at the potential expense of prolonging the life-time of temporary memory However, many of the optimization techniques described in this will effect more significant trade-offs between code chapter execution time, and temporary storage. As a result, similarity functions are to provide the necessary data in terms of which the desirability of initiating a particular optimization can be measured. The feasibility of the optimizations exposed by similarity will be described in terms of the

primitives developed in Chapter II.

This aspect of a similarity function is illustrated by the following example:



Assume that the previous optimization strategy which moved part of the evaluation to a point after the merge is precluded by the fact that es and end produce side-effects on constituents of the assignment to A and so block the forward motion of the assignments within their respective linear blocks. A possible optimization of this example consists of replacing the control expression above with

The jump to and return from the common evaluation point of the expression $A \leftarrow e_2 + e_3 * (e_4 + T)$ incurs an overhead cost associated with the execution of the additional control. This overhead did not occur in the preceding example since the common code sequence was entered after the merge. The decision to invoke this optimization must be made in terms of the trade-off between (1) the amount of code saved by the common code sequence, and (2) the space and time overhead incurred by the introduction of additional code for control. The measurements involving code size are computable at compile time. Measurements involving execution time can at best only be estimated at compile time with assumptions relating to factors such as depth of loop nesting and equi-probable selection of parallel branches in forked control environments. If the similarity function is to

be a useful tool for describing optimization strategies such as those listed above, then it must provide the information required to evaluate these trade-offs.

The two preceding examples involved optimization techniques that altered the control flow of the original program. The second example models the standard programming construct of a call to a single parameter procedure. This latter observation opens up a whole spectrum of applications for the similarity concept. The fact that the two assignments to A are on parallel branches of a forked construct is not essential to the feasibility of the optimization strategy. The effectiveness of invoking the optimization technique is measured in terms of the trade-off between the cost of the parameter mechanisms and calling sequence overhead and the cost of the parameterized code sequences which the "almost identical" expressions can share.

EXAMPLES OF SOME SIMILARITY FUNCTIONS

Chapter II presented a minimal set of requirements to be satisfied by a similarity function. They reflect the notion that members of this family of functions are to "measure" the degree of identity of the superstructure of pairs of expressions.

 $\boldsymbol{\sigma}$ is a similarity function only if

(1) σ : ExE \rightarrow [0, ∞),

(2)
$$\sigma(e_1,e_2)=\emptyset$$
 iff $e_1\cong e_2$, and

(3)
$$\sigma(e_1,e_2) = \sigma(e_2,e_1)$$
 for all $e_1,e_2 \in E$. (symmetric)

These requirements alone, however, do not suffice to convey the notion of a measurement of "almost-congruence". For example the function:

$$C(e,e') = \begin{cases} \emptyset & \text{if } e \cong e' \\ 1 & \text{otherwise} \end{cases}$$

satisfies requirements (1)-(3) but conveys precisely the same information as the "sa" relation. A similarity function must provide a more selective measurement.

A first approximation to a similarity function is provided by the function F defined below. F does a coordinated tree walk returning from each corresponding pair of nodes which are not congruent with a value of 1. The following algorithm't gives a precise description of F.

[†] These algorithms are presented in pseudo-Bliss. Their translation to "true" Bliss would require specification of data formats to a level of detail exceeding our present needs.

end;

```
routine F(e,e')=
            ! N is the set of non-terminals of F.
            ! T is the set of terminals, L is the set
            ! of literals and I the set of names.
        begin local S:
            if e ∈ N xor e' ∈ N then return 1;
            if e ∈ N then
                 begin
                      if e[operator] ≠ e'[operator]
                          then return 1;
                     if e[# of operands] ≠ e'[# of operands]
                          then return 1;
                     S←Ø:
                     incr I from 1 to e[# of operands] do
                          S .- S+F(e[operand;], e'[operand;]);
                     return .S:
                end:
            if e ∈ L xor e' ∈ L then return 1;
            if e ∈ L
                then return literalvalue(e) # literalvalue(e*):
           return name(e) ≠ name(e')
```

Example

Let e_0 : .E+.B\$.D, e_1 : .A+.B\$.C, e_2 : .A+.B\$.(C+.E), \ni_3 : .A+.B\$.(D+.E). Then $F(e_0,e_1)=2$, $F(e_1,e_2)=F(e_1,e_3)=F(e_2,e_3)=1$.

In effect, F provides a count of the number of dissimilar nodes in a pair of expressions. It does not, however, provide a very selective measure since it does not distinguish between the pairs (e₁,e₂) and (e₂,e₃). The expressions e₂ and e₃ are more alike in some intuitive sense because their dissimilarities occur at a lower level in the tree. A function which could distinguish between such pairs would be preferable since differences at a greater depth correspond to longer identical code sequences.

The following function G incorporates the weighting factor of tree depth by a slight modification of F. The difference between F and G occurs at the point of the recursive call where the reciprocal of the number of operands is inserted as a multiplicative factor. The procedure for G is:

```
routine G(e,e')=
        begin local S;
             if e ∈ N xor e' ∈ N then return 1;
             if e ∈ N then
                 begin
                      if e[operator] ≠ e'[operator]
                          then return 1;
                      if e[# of operands] ≠ e'[# of operands]
                          then return 1;
                      S←Ø:
                     incr I from I to e[# of operands] do
                          S←.S+(1/e[# of operands])*
                              G(e[operand;], e'[operand;]);
                      return .S:
                 end:
             if e ∈ L xor e' ∈ L then return 1;
            if e € L
                 then return literalvalue(e) # literalvalue(e');
             return name(e) ≠ name(e')
        end:
```

The function, G, produces a finer measure on pairs of expressions than F. For example, let e_1 , e_2 and e_3 be as defined in the preceding example. $G(e_1,e_2)=0.25$ and $G(e_2,e_3)=0.125$ whereas $F(e_1,e_2)=F(e_2,e_3)=1.0$. However, if we define e: (.A+1)*(.A+2)*(.A+3), e': (.B+1)*(.B+2)*(.B+3), and e'': (.A+1)*(.B+2)*(.C+3), then $G(e,e^i)=G(e^i,e^{ii})=G(e,e^{ii})=0.5$. Hence $G(e,e^i)=G(e,e^i)=G(e,e^i)=0.5$.

The final similarity function presented here is the one that has been implemented in the optimization pass which produces the examples in Chapter IV. SIGMA initializes the variables NPARMS to zero and COSTAV to the estimated object code size of the expression e. Whenever the recursive subroutine S encounters a pair of dissimilar subexpressions of e and e' it calls the subroutine TRYPARMS. The subroutine TRYPARMS determines if a new parameter must be created. If so, it increments NPARMS by one and decreases COSTSAV by the estimated object code size of the parameter subexpression of e. When control returns to SIGMA from S, the variable NPARMS contains the number of parameters necessary to evaluate the pair e, e' by a common code sequence and COSTSAV contains an estimate of the size of the object code sequence sharable by the expressions.

routine SIGMA(e,e')=

! The subroutine S does a coordinated tree walk on the ! expressions e and e' setting the variables NPARMS to ! the number of parameters and COSTSAV to the amount ! of code saved by the shared code sequences. The ! subroutine TRYPARMS (not defined here) increments ! NPARMS and decrements COSTSAV by e[cost] if a new ! parameter must be created. e[cost] is the amount ! of code necessary to evaluate the entire expression e. ! e[count] is the number of formally identical ! instances of this expression.

begin own NPARMS, COSTSAV, M;

```
routine S(e,e')=
        begin
             if e ∈ N xor e' ∈ N
                 then return TRYPARMS(e,e');
             if e ∈ N then
                 begin
                     if e[operator] ≠ e'[operator]
                          then return TRYPARMS(e,e');
                     if e[# of operands] # e'[# of operands]
                          then return TRYPARMS(e,e');
                     if e ≅ e' then return;
                     incr I from 1 to e[# of operands] do
                          S(e[operand, ],e'[operand, ]):
                     return
                 end:
             if e € L and e' € L
                 then return
                     if literalvalue(e) ≠ literalvalue(e')
                          then TRYPARMS(e,e');
             if e € land e' € l
                 then return
                     if name(e) # name(e')
                          then TRYPARMS(e,e');
             TRYPARMS(e.e')
        end:
    if e ≅ e' then return Ø;
    NPARMS←Ø; COSTSAV←e[cost];
    S(e,e');
    M←e[count] + e'[count];
    (.M*.NPARMS+.M+1)/((.M-1)*.COSTSAV)
end;
```

The final expression in the body of SIGMA requires some explanation.

The numerator is the estimated cost in code size of the overhead required to set up parameters (.Me.NPARMS), call (+.M), and return (+1) from a

similarity-created subroutine. The denominator is the amount of code saved by replacing M-1 of the expressions with calls to a common sequence of code. Hence if SIGMA(e,e')<1, then code size will be reduced by implementing e and e' as calls on a common subroutine.

The application of the similarity function SIGMA (more precisely its subroutine S) partitions an expression into a body and a collection of parameter expressions. In subsequent discussions, body(e) refers to the expression resulting from the removal of the parameter nodes in e, and parms(e) refers to the set of sub-expressions identified by S as parameters of e.

Example

Define: e₁: (.A+1)*(.A+2)*(.A+3), e₂: (.B+1)*(.B+2)*(.B+3), e₃: (.A+1)*(.B+2)*(.C+3), e₄: .A+.B*(.C+.E), e₅: .A+.B*(.D+.E), e₆: .A+.B, e₇: .A+.C

The following table shows the values returned from F, G, and SIGMA.

	F		SIGMA			
e ₁ ,e ₂	3.0	Ø.5	0.625 = (2*1+2+1)/8			
e ₁ ,e ₃	3.0	Ø.5	1.125 = (2+3+2+1)/8			
e4,e5	1.0	0.125	1.5 = (2*1+2+1)/4			
e ₆ ,e ₇	1.0	0.5	$2.5 = (2 \div 1 + 2 + 1)/2$			

It must be emphasized that we have presented an example of a particular similarity function that has produced extremely interesting results in our optimization pass. There are a variety of such functions each sharing common basic characteristics with SIGMA. Indeed this

particular similarity function ignores the execution time overhead resulting from introducing subroutine linkages and so identifies those optimizations that minimize object code size as "desirable" without regard to their effect on execution time.

Throughout the remainder of this chapter, the existence of a similarity function, σ , whose essential characteristics are mirrored by SIGMA and its subroutine S will be assumed. The following sections will present a collection of optimization techniques defined in terms of similarity and the primitives defined in Chapter II.

CONVERTING EXPRESSIONS TO SUBROUTINES

A programmer selects macros and procedures to define in his program on the basis of logically coherent units of computation. Macros (expanded in line) save time by avoiding execution time linkage and parameter passing mechanisms at the expense of increasing object code size. Closed procedures, on the other hand, reduce object code size at the expense of run-time overhead. The decision to choose a macro over a procedure or vice versa is typically made on the basis of some rough and usually intuitive estimate of the ratio of the object code size to the frequency of occurence.

An optimization strategy described in terms of similarity eliminates

this decision for the programmer by expanding the simple (i,e. non-recursive) procedures in line. The decision to close some of these procedures or portions of them is made on the basis of information collected by a similarity function. This process, that identifies expressions to be implemented as closed subroutines, operates only on the form of the program. As a result it can identify computationally coherent sequences which do not possess a logical coherence that would lead to their identification as a macro or procedure by the programmer. The examples in Chapter IV demonstrate that these situations occur in real programs!

Similarity can be used to identify those expressions which occur sufficiently often that their implementation as subroutines will reduce object code size. The similarity function SiGMA returns a value which is the ratio of the overhead to amount of code saved by creating a subroutine out of a pair of expressions. If that ratio is less than 1, then a savings in code size results.

As we mentioned above, the value returned from SIGMA(e,e') indicates whether a subroutine creation is desirable, however it does not imply that such a creation is also feasible. Consider the example of an expression e that is to be implemented as a subroutine with a single parameter. Furthermore assume that for one of the calls on e the actual parameter expression contains .X as an operand. Finally assume that the subroutine implementation adopts a call-by-value convention for parameters. Thus, the value of X will be accessed during the parameter evaluation prior to

evaluation of the expression e. If the expression e alters the value of X prior to the original evaluation point of the parameter expression, then the data flow semantics for e have been violated. Furthermore since the parameter expression can appear within a loop contained in e, it is not sufficient that no re-evaluation of X precede the parameter expression. This set of observations can be summarized as follows:

A subroutine creation from the expression e and e' is feasible if p \in prolog (β) Ω epilog(β) \forall p \in parms(e), p' \in prolog(β ') Ω epilog(β ') \forall p' \in parms(e'), where $\beta = \beta$ (cover(e) and β ' = β '(cover(e').

The criterion that SIGMA(e,e')<1 is sufficient to guarantee that the subroutine implementation of e and e' will reduce code size. It is quite reasonable to define a controlling heuristic that weighs the amount of code saved against the storage required for parameters (especially if they are passed in registers) and some expected value of increased execution time. This observation argues for a function DELTA which is SIGMA dependent and encodes the heuristics to be applied in deciding the desirability of implementing a set of expressions as a subroutine. Hence the decision to evoke these optimizations will be made by a predicate of the form: $\sigma(e,e')$ < $\delta(e,e')$. Logically the function δ is defined in terms of the expressions e and e'. However, in an implementation of δ , one expects the subroutine DELTA to share information collected by SIGMA. In particular DELTA should have access to the own variables NPARMS and COSTSAV. A straightforward extension of the notion of strong similarity makes the dependence of δ on the expressions e and e' explicit: $e \approx e'$ iff $\sigma(e,e') < \delta(e,e')$. The

examples in Chapter IV demonstrate the results that occur when δ is set to a constant value of 1.0. We will refer to this optimization technique which creates subroutines from sets of strongly similar expressions as the strong similarity subroutine optimization (\$3 optimization). Throughout the remainder of the chapter, we will fix the interpretation of δ to be 1.0 and as a result $e \simeq e^t$ iff $\sigma(e,e^t) < \delta = 1.0$.

PARTIAL POST-EVALUATION IN FORKS

The S3 optimization will generally use subroutine call and return instructions in its implementation. The next few sections point out cases that simplify the linkage mechanism.

Reconsider an example presented earlier in the chapter

$$\begin{array}{c} \underline{\text{if}} \ e_0 \\ \underline{\text{then}} \ (e_1; ...; e_K; \ A \leftarrow .B + .C \div .D) \\ \underline{\text{else}} \ (e_{K+1}; ...; e_n; \ A \leftarrow .E + .C \div .E). \end{array}$$

The two assignments to A are strongly similar making it feasible to apply an S3 optimization to them. However, the optimization:

$$\begin{array}{c} (\underline{if}\ e_0\\ \underline{then}\ (e_1;\ ...\ ;\ e_k;\ T\leftarrow D)\\ \underline{else}\ (e_{k+1};\ ...\ ;\ e_n;\ T\leftarrow E);\ A\leftarrow B+.C+.T). \end{array}$$

avoids a subroutine mechanism. The following general description applies to optimizations of this form:

Given an n-way branching environment

$$\beta_1$$
 β_2 ... β_n

a partial 'post-evaluation of the strongly similar expressions $e_1 \in \beta_1, \dots, e_n \in \beta_n$ is feasible if body(e₁) \in postlog(β_1), ... , body(e_n) \in postlog(β_n) and $p \in$ prolog($\beta(cover(e_1))$) \cap pellog($\beta(cover(e_1))$) $\forall p \in parms(e_1), 1 \le i \le n$.

This optimization is accomplished without adding additional linkage mechanism and so saves space without increasing execution time.

WASP-WAISTING -- REVISITED

In Chapter II a brief reference was made to an optimization we called "wasp-waisting". A representative example is

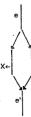
$$\begin{array}{c} \underline{\text{if }} \ e_0 \\ \underline{\text{then}} \ (e_1; ...; e_i; ... ; e_j) \\ \underline{\text{else}} \ (e_{i+1}; ... ; e_k; ... ; e_n) \end{array}$$

where $e_i \cong e_k$. The optimization strategy for this example consisted of replacing the expressions e_i and e_k with calls to a common subroutine. The S^3 optimization technique extends this strategy to cases in which the expressions e_i and e_k are not formally identical but only strongly similar. The feasibility requirements for a wasp-waisting optimization are identical

to the requirements for an S3 optimization. The difference between the two optimizations lies in the possibility of implementing the subroutine call by a simple branch instruction and the return by retesting the selector.

GENERATING CONDITIONAL SUBROUTINES

Consider the following example



where $e \approx e'$ and e' is an expression involving X. Compile time data flow analysis clearly indicates that e and e' are not redundant expressions because of the potential assignment to X. At run time, on the other hand, whenever control passes down the right hand branch, the post-merge evaluation of e' is unnecessary.

An optimization strategy consists of replacing e and e' by calls on a subroutine which conditionally executes depending on a boolean value. The form of the subroutine is

if <boolean> ihen (<subr-body>; <boolean>+false).

The boolean is set to true initially and reset to true at the point of the store into X or anywhere on the left branch of the fork. This optimization technique saves space as does any S3 optimization. It also saves time presuming that the time to evaluate the subroutine body exceeds the time involved in setting and testing the boolean.

GENERATING LOOPS

The optimizations described by similarity to this point have involved the introduction of additional branches and subroutine calls. This section will investigate two optimization strategies described using similarity which introduce loops into a program.

Consider a compound expression $e: (e_1; e_2; ...; e_n)$ in which $e_i \simeq e_j, \ 1 \le i,j \le n$ and (for simplicity) assume the set parms (e_i) is a singleton $\{p_i\}, \ 1 \le i \le n$.

Case 1:

Independent of the relationship between the corresponding parameters of the expressions, this compound expression can be implemented by a control environment which models the Algol <u>for</u>

 $\underline{\text{for}} \mid l := p_1, p_2, \dots, p_n \underline{\text{do}} e^{\text{"} < l}$ where $e^{\text{"} < l} > \text{is body}(e_i)$ in which $parm(e_i) = l$. Case 2:

If in addition the parameter expressions are such that $p_{i-1}\text{-}p_i = \Delta p, \ 1 \leq i < n \ \text{ and } \Delta p \in \text{prolog}(\beta|e) \ \Omega \ \text{epilog}(\beta|e) \ \text{then } \ \text{the compound expression } e \ \text{can be implemented by}$

incr I from p₁ to p_n by Δp do e"<1>

where e"<1> is as described in case 1.

The restriction to single parameter subroutines can be removed by updating a set of control variables, one for each parameter, on each iteration of the loop. Both examples reduce object code size and replace the subroutine linkage mechanism of the S³ optimizations with loop control. In addition case 2 reduces both time and space costs by incrementally computing successive parameters.

The pair of optimization strategies relates quite closely to our discussion of strength reduction in Chapter II. In particular we sought a technique for discovering a relation F such that given a pair of expressions e and e¹, $F(e,\Delta e)=e^t$. Both cases 1 and 2 provide solutions. The relation F is precisely the loop body expression e" and the parameter Δe is the loop variable I. The fact that $e^t \simeq e$ guarantees that the size of the object code to compute $F(e,\Delta e)$ is less that that required to evaluate e' in the ususal manner. Case 2 also demonstrates the discovery of an inductive relationship among the expressions e_1 , e_2 , ..., e_n without assumptions on the form of the expressions. In particular no restriction to polynomial expressions is required.

SIMILARITY AND ITERATIVE TECHNIQUES

The next optimization described in terms of similarity arises often in algorithms concerned with various forms of iterative analysis. A simple example motivates the usefulness of the optimization strategy.

The following algorithm accumulates in S the trapezoidal approximation to the definite integral of F over the interval $[X_0,X_n]$:†

$$\frac{incr}{S \leftarrow S} + \frac{from}{(F(.I-\Delta X) + F(.I))/2)} \times \frac{\Delta X}{\Delta X} \times \frac{do}{\Delta X}$$

The important item to note here is that on the k-th iteration of the loop the value of $F(.I-\Delta X)$ is precisely the same as the value of F(.I) on the (k-1)-st iteration. Recognizing this relationship between $F(.I-\Delta X)$ and F(.I), an optimization strategy that requires only one evaluation of F per iteration is given by:

```
\begin{array}{ll} \underline{if} & X_0 + \Delta X & \underline{leq} & X_n & \underline{then} & OLDF \leftarrow F(X_0);\\ \underline{incr} & 1 & \underline{from} & X_0 + \Delta X & \underline{to} & X_n & \underline{by} & \Delta X & \underline{do} \\ \underline{begin} & & NEWF \leftarrow F(.1);\\ & S \leftarrow S \leftarrow ((.OLDF + .NEWF)/2) \neq \Delta X;\\ & OLDF \leftarrow .NEWF \\ \underline{end}. \end{array}
```

A description of the expressions in a loop body for which this optimization is feasible is given by:

[†] The loop models the "calculus-text" description of the trapezoidal, rule, even though a numerical analyst would not program it in this form.

Given a loop incr I from e_0 to e_1 by e_2 do e_3 where f and g are subexpressions of e_3 , $f \simeq g_3$ let parms($\beta = [p]$ and parms($g) = [p + e_2]$. Then it is feasible to eliminate the evaluation of f on each iteration of the loop (and replace it by the old value of g) if body($f) \in \text{prolog}(\beta_3)$ $f) = \text{pilog}(\beta_3)$ and body $f) \in \text{prolog}(\beta_3)$ $f) = \text{pilog}(\beta_3)$ and $f) = \text{pilog}(\beta_3)$ $f) = \text{pilog}(\beta_3)$ $f) = \text{pilog}(\beta_3)$

The restriction that body(f) and body(g) be loop invariant is equivalent to stating that on any two iterations of the loop the evaluations of f(c) and g(c) produce the same value for a fixed parameter c.

SUMMARY

In the preceding sections the similarity function has been used to decribe a variety of apparently unrelated optimization strategies. This fact reflects its usefulness as a unifying primitive which can be employed in describing a wide range of concepts. Indeed this property may be sufficient justification in itself for proposing the similarity notion.

However, the ensuing chapter presents a strong case that similarity has very practical application in an optimizing compiler. The reductions in code size that result from application of the S3 optimization alone are remarkable. In addition the S3 optimizations demonstrate interesting results in identifying computationally coherent expressions from analysis of a program's form. Sometimes these computationally coherent expressions correspond to those which the programmer considered logically coherent and

CHAPTER IV

Chapters II and III propose a number of optimizations. This chapter discusses the relative significance of some of these optimizations in some specific cases. A program is described which implements both the optimizations described in Chapter II and the S3 optimization (using the particular similarity function, SIGMA, described in Chapter III). The chapter is subdivided into two parts: (1) a description of the program and the form of its output and (2) a discussion of a set of examples which show the effect of the optimizations.

The purpose of the program is the evaluation of the effectiveness of the \$3 optimization as compared with the classical optimizations of Chapter II. Therefore, it was not to our purpose to construct a complete compiler. However, since the Bliss-10 compiler already accepts Bliss syntax and produces PDP-10 machine code, we have chosen the PDP-10 as the target machine and will demonstrate shortly that the estimates produced by our program correspond to the actual number of machine language instructions produced by Bliss-10. To enable the comparison between the classical optimizations and \$3\$, the program is constructed so that programs may be compiled with various subsets of the optimizations enabled.

KATE

The program KATE† is a translator from Bliss to a three-address code. KATE may be thought of as four modules each of which works on a representation of the program and global information prepared by other modules. The first module, LEXSYN, performs lexical and syntactic analysis on the source text, builds a symbol table and produces a tree representation of the program as described in Chapter I. The second module, FLOW, implements the optimizations described in Chapter II except for strength reduction. Ommision of strength reduction does not effect the comparison between the classical optimizations and S3 since it has relatively little effect on object code size. The third module, S3, implements the S3 optimization described in Chapter III. The fourth module, CODE, produces a three address code and an estimate of the number of PDP-10 instructions that would result if the three address code were translated into real machine code. The following diagram illustrates the possible paths which KATE can follow in translating source text to three address code.

[†] For those who feel that acronyms require interpretations, we suggest Algorithmic Translating Engine. The K, of course, is silent.

source
$$\rightarrow$$
 LEXSYN \rightarrow FLOW \rightarrow CODE \rightarrow 3-address code LEXSYN \rightarrow FLOW \rightarrow FLOW \rightarrow S3+

Although the FLOW module can be thought of as a separate pass over the representation produced by LEXSYN, in fact, it processes the tree from the inside out while the tree is being built by LEXSYN. The FLOW module is invoked by LEXSYN at the completion of each linear block to build the prolog, epilog, and postlog sets. As syntactic analysis is completed for each control environment, flow is called to invoke the various optimization strategies. In particular a node representing a forked control expression points to the α and ω sets for the expression and each node representing a looping control expressions points to the χ (loop invariant expressions) and ρ (cyclic re-evaluations) set for that expression. S³, on the other hand, makes a completely separate pass since it must have information on all occurrences of strongly similar expressions to make its decisions. LEXSYN, FLOW, and S³ implement the primitives developed in Chapters II and III. However, CODE requires more detailed explanation of its output to facilitate understanding of the examples.

CODE

The CODE module translates the tree representation of the program into

a three address code. The machine code operations which are used in CODE were selected to facilitate an accurate estimate of the number of PDP-10[P70] machine language instructions that would result from the three address code. Again, the PDP-10 was chosen because the Bliss-10 compiler enabled us to verify the accuracy of the estimates made by CODE.

The three address code is formatted as:

operator operand₁, operand₂, operand₃.

Each operator has a fixed number (0,1,2,3) of operands. The operand of an instruction can be:

- (1) a name -- e.g. X
- (2) the value pointed to by a name -- e.g. .X
- (3) a level of indirect on on (2) -- e.g. ..X
- (4) a constant
- (5) a label.

The following table lists the machine code operations and describes their semantics. In general e₁ and e₂ are the operands of the opcode and e₃ is the result returned to the parent node of the subnode which produced the result.

			A.
opcode	operands	semantics	
ADD		e ₃ -e ₁ +e ₂	
SUB	e ₁ ,e ₂ ,e ₃	e ₃ ←e ₁ -e ₂	
MUL	e ₁ ,e ₂ ,e ₃	e3←e1*e2	
LTSH	e ₁ ,e ₂ ,e ₃	e ₃ ←e ₁ 1e ₂	
RTSH	e ₁ ,e ₂ ,e ₃	e ₃ ←e ₁ ↑(-e ₂)	
DIV	e ₁ ,e ₂ ,e ₃	e ₃ ←e ₁ /e ₂	
MOD	e ₁ ,e ₂ ,e ₃	e ₃ ←e ₁ mod e ₂	
GTR	e1,e2,e3	e ₃ ←e ₁ >e ₂	
LEQ	e ₁ ,e ₂ ,e ₃	e ₃ ←e ₁ ≤e ₂	
LSS	e1,e2,e3	e ₃ ←e ₁ <e<sub>2</e<sub>	
GEQ	e ₁ ,e ₂ ,e ₃	e ₃ ←e ₁ ≥e ₂	
EQL	e ₁ ,e ₂ ,e ₃	e ₃ ←e ₁ =e ₂	
NEQ	e1,e2,e3	e ₃ ←e ₁ ≠e ₂	
AND	e ₁ ,e ₂ ,e ₃	e ₃ ←e ₁ and e ₂	
ANDCR	e1,e2,e3	e ₃ ←e ₁ and not e ₂	
ANDCL	e1,e2,e3	e ₃ ← <u>not</u> e ₁ and e ₂	
ANDCB	e1,e2,e3	e ₃ ← <u>not</u> e ₁ and not e ₂	
OR	e1,e2,e3	e ₃ ←e ₁ <u>or</u> e ₂	
ORCR	e1,e2,e3	e ₃ ←e ₁ <u>or not</u> e ₂	
ORCL	e1,e2,e3	e ₃ ←not e ₁ or e ₂	
ORCB	e1,e2,e3	e ₃ ← <u>not</u> e ₁ <u>or not</u> e ₂	
LD ·	e ₁ ,e ₂ ,e ₃	e ₁ ←e ₂ (e ₃ is value of exp.)	
LDN	e1,e2,e3	e ₁ ←-e ₂ (")	
LDC	e1,e2,e3	e ₁ ← <u>not</u> e ₂ (")	
XCT	e ₁ ,e ₂	execute inst. at e ₁ +e ₂	
PARM	e ₁	set up parameter eı	Jassumes stack
DPARM	e ₁	deallocate e ₁ parameters	(discipline
CALL	e ₁ ,e ₂	save PC+1; PC←e1; value in e	2
RTRN		PC←saved value	
BR	e ₁	PC←e ₁	
BRT	e ₁ ,e ₂	if e₁ then PC←e₁	
BRF	e ₁ ,e ₂	if not e₁ then PC←e2	
INC	e ₁	e1 - e1 + 1	
DEC	eı	e1←e1-1	
SSCAL	e ₁ ,e ₂ ,e ₃	save e2; PC←e1; value in e3	

An example of output from KATE demonstrates the use of these operations.

<u>begin own</u> I,V[10],X,Y,Z,A,B,C,D,F; X←-,Y-,Z; V[2*.I]←-,Y-,Z; Z←,A*.B+,C*.D; F(.Z) <u>end eludom</u>

ADD	Y, Z, T ₁	2*
LDN	X,.T ₁ ,T ₁	1 *
MUL	2, I, T ₂	2=
ADD	$V_{\nu}T_{2\nu}T_{2\nu}$	Ø÷
LDN	$.T_{2}, T_{1},T_{1}$	1 *
MUL	.A,.B,.T ₁	. 2≠
MUL	.C,.D,.T2	2*
ADD .	$.T_{1},.T_{2},.T_{1}$	1 *
LD	Z_1, T_1, T_1	
PARM	.Z	1*
CALL	F _s .T ₀	1 *
DPARM	1	1 *

TOTAL COST= 15

The following points should be noted:

- (1) The code generators are table driven. They attempt to do peephole optimization on a very local level. For example, note that the expression -.X-Y was converted to -(X+Y).
- (2) CODE allocates temporary storage (T_1 and T_2) in a straightforward manner. If a temporary location is allocated to a redundant expression, it remains reserved for the value of that expression until the last occurrence of the use of that value. In the example above the product $2 \cdot 1$ was formed in T_2 since the last use of the value in T_1 followed the index computation.
- (3) Note that there is only one instruction for transferring the value stored in one memory location to another. In particular a LD instruction can correspond to (a) loading a temporary -- LD T_{0} , A, (b) storing a temporary into a user-defined memory location -- LD A, T_{0} , or (c)

transferring the contents of one memory location to another -- LD A,B.

(4) The column of numbers to the right contains estimates of the number of PDP-10 instructions required by each operation. An actual PDP-10 code sequence for this example is:

```
T_1,Y
MOVE
ADD
              T_1,Z
MOVNM
              T_1,X
MOVE
              T_2,I
IMULI
              T2,2
                       (or: ASH
                                     T_{2},1)
              T_1,V(T_2)
MOVNM
MOVE
              T<sub>1</sub>,A
              T<sub>1</sub>,B
IMUL
MOVE
              T2,C
!MUL
              T<sub>2</sub>,D
ADD
              T_1,T_2
MOVEM
              T<sub>1</sub>,Z
PUSH
              $S.Z
                       $$ points to the stack
PUSHJ
              $S.F
              $5,[10000001]
SUB
```

In particular notice that KATE estimates ϑ as the cost of the operation $V+T_2$ since the addition can be accomplished by indexing.

(5) The indentation exhibiting the columns of asterisks indicates the nesting of linear blocks and the number at the base of a column is the cumulative total of the code size for that linear block. This facilitates comparison of the number of instructions in critical regions such as inner loops.

VALIDATION OF KATE'S ESTIMATES

The estimates of object code size are generated on an instruction by

instruction basis. Corresponding to each machine code operation produced by KATE there is a 12x12 table. An index into the table is computed by analyzing each operand into one of twelve states:

L is a literal (absolute value greater than 1), N is a user defined storage location, T is a compiler defined temporary (whose contents may be destroyed by the execution of the instruction), and T' is a temporary whose contents must be preserved.

In order to demonstrate that the numbers produced by KATE are in fact reasonable when applied to sequences of code, a comparison was made between the estimates produced by KATE and actual PDP-10 machine code produced by Bliss-10. Both compilers were run with all optimization turned off. This was done since even though the two compilers apply different sets of optimizations, they both produce straightforward, simple machine code with all optimizations turned off. We have selected two examples (to be examined in more detail for other purposes later in the chapter) to exemplify the results. The first example is a large sub-program taken from the Bliss-10 compiler itself. Bliss-10 produces 983 PDP-10 instructions. The estimate produced by KATE is 979 instructions. The difference is less than 0.5%

A second example, an implementation of the quadratic formula, is small enough to be reproduced in its entirety. The source text is the following:

The output on the left column of the next page is produced by KATE; the output on the right is produced by Bliss-10. In the Bliss-10 output: A = -4(\$F), B = -3(\$F), and C = -2(\$F).

	KATE			ı		Dlic	- 10
	NAIL			1		DIIS	<u>s-10</u>
POOT:				I		JSP	
HUL	.YYT\$1	2		1		MOVE	12ENT.0 043(\$F)
MUL	1X1\$2	2 *		1		IMUL	04,-3(\$F)
HUL	.1522152	1 =		i		MOVE	054(\$F)
SUB	. T\$1 T\$2 T\$1	1 -		Į.		ASH	05.2
LSS	.751.0751 .751.L51	0 .		ł		IMUL SUB	05,-2(\$F)
LD	ERROR.1.1	2		ł		JUMPGE	04.5 04.12020
BR	L\$2	1 .		i		HOVEI	\$V.1
		3		ı		HOVEH	SV.ERROR
L\$1: MUL	.YYT\$1	2 .		I	L2020	JRST	\$5.L1536
MUL	4XT\$2	2 =		1	L2020:	MOVE	07,-3(\$F) 07,-3(\$F)
MUL	.7\$2521.	i .		1		MOVE	10,-4(\$F)
SUB	.751752751	1 .				ASH	10.2
EQL	. 151.0, . 151	. 6 .				INUL	102(SF)
BRF HUL	.T\$1.L\$3 2XT\$1	2 =		1		SUB	07,10
DIV	.YTS1,TS1	2		1		HOVE	07.L2406 124(\$F)
LDN	R1,.T\$1,T\$1	i :				ASH	12.1
MUL	2xT\$1	2 .		i		MOVN	053(\$F)
DIV	.YT\$1T\$1	2 ×		i		IDIA	05.12
EDN BR	RZTS1TS1 LS4	1 *		l		HOVEH	05.R1
Dr.	L91	1 11		í		HDVE ASH	\$V4(\$F) \$V.1
L\$3:		••				HOUN	04,-3(\$F)
MUL	2XT\$1	2 ■		ì		IDIV	04.3
DIV	.Y,.T\$1,T\$1	2 .				HOVEH	04.R2
MUL MUL	.Y,.Y,.T\$2 4,.X,.T\$3	2 .			L2406:	JRST HOVE	\$S.L1536
MUL	.7\$321\$3	í :		l .	LZ4061	INUL	\$V,-3(\$F) \$V3(\$F)
SUB	.1\$27\$31\$2	i i		ľ		HOVE	12,-1(\$F)
PARM	.1\$2	1 ≪				ASH	12.2
DPARM	SORT TSO	1 .				IMUL	122(\$F)
MUL	2XT\$2	2 *		i		SUB PUSH	\$V.12 \$5.3
DIA	.150,.152,.750	1 .		1		PUSHJ	\$S.SQRT
SUB	.750751750	j .				SUB	\$5,1000001,,00000011
LD	R1TS0TS0	1 =		l		HOVE	054(\$F)
DIV	2XT\$1 .YT\$1T\$1	2 2		1		ASH	05.1
HUL	.YYT\$2	2 .		l		HOUN	063(\$F) 06.5
MUL	4XT\$3	2 .				MOVE	04,-4(\$F)
HUL	.1\$321\$3	1 =				ASH	04.1
SUB PARM	.T\$2,.T\$3,.T\$2 .T\$2	! :				IDIV	\$V.4
CALL	SQRTTS0	1 :				ADD HOVEH	\$V.6 \$V.R1
DPARM	1	i :				HOVE	\$V,-3(\$F)
HUL	2XT\$2	2 *				IHUL	\$V,-3(\$F)
DIV	.150152150	1 =				HOVE	10,-4(\$F)
ADD LDN	.T\$1T\$0,~.T\$1 R2T\$1,T\$1	1 :				ASH	10.2
Lon	K21.1917-1191	1 •			ί.	1MUL SUB	10,-2(\$F) \$V.10
L\$4:		30			•	PUSH	\$5.3
		54				PUSHJ	\$5.SQRT
L\$2:						SUB	\$5.10000010000011
LD RTPN	T\$0.0T\$0	1:				MOVE	11,-4(\$F)
KIRN		66				MOUN	11.1 123(\$F)
		•				IDIV	12,-3(%)
						MOVE	054(\$F)
TOT	AL COST= 66					ASH	95.1
						1DIV SUB	\$0.5
						HOVNH	\$V.12 \$V.R2
			- 1		L1536:	SETZ	\$0.0
			- 1			JRST	\$SEXT.0
			- 1				
			- 1		MODULE	LENGTH =	740
			i		HODOLE	LENGTH D	27-0

The additional instruction (JSP 12,ENT.0) in Bliss-10 executes routine entry code. These examples demonstrate that the estimates of object code size produced by KATE are indeed reliable predictions of the actual number of PDP-10 machine language instructions that would be generated from the three-address code.

The remainder of Chapter IV discusses three examples which contrast the effect of the classical optimizations and the S3 optimization introduced in Chapter III. The examples demonstrate the potential of the S3 optimization for producing significant reductions in object code size. KATE was run in three modes on the examples: (1) NOOPT: no optimizations, (2) ALLBUTSIM: S3 by-passed, (3) ALLOPT: S3 included.

OUADRATIC FORMULA

The first example involves three implementations of a program to evaluate the quadratic formula. The main routine, ROOT, is identical in all three implementations. The difference occurs in the evaluation of the square root.

EXAMPLES QUADRATIC FORMULA

R3:

```
PEGIN
   R1:
                                                 HOLEU
                                                   POSPODI(A.B.C)=(-B)/(2+A)+SQRT(DISC(A.B.C))/(2+A)$.
                                                    NEGRODI(A.B.C:=(-B)/(Z+0)-SQRT(DISC:0.B.C))/(Z+0)$.
SORT
              is
                              subroutine
                                                   DISC(A.B.C)+B+8-4+A+C$:
                                                  FORMARD SORT:
                                                  GLOBAL EPROP.R1.R2:
implementation of the sequence
                                                  ROUTINE POOT(X,Y,Z)=
                                                    BEGIN
                                                      IF DISC(.X..Y..Z) LSS @ THEN EPPCP-1 ELSE
\{(x_n^2+A)/(2x_n)\}
                                    which
                                                      IF DISC(.X..Y..Z) EQL 9 THEN
(R1+-.Y/(2*.X)) P2+-.Y/(2*.X))
                                                      ELSE (R1-POSROOT(.x..Y..2):R2-NEGPCCT(.x..Y..2));
converges to square root of A.
                                                    END:
                                                  ROUTINE SORT(X)*
                                                    BEGIN
                                                      LOCAL XI.XJ: GLOBAL EPSILON: MACPO INFINITY=#7777775:
(Newton's method).
                                                      XI-.X: XJ-INFINITY:
                                                      WHILE (.XJ-.XI) GTP .EPSILON
                                                        DO (XI+.XJ: XJ+(.XI*.XI+.X)/(2*.XI));
                                                      .x.
                                                    END:
                                                CND ELIDON
   R2:
                                                BEGIN
                                                    POSROOT(A.B.C)=(-B)/(Z*A)*SQRT(DISC(A.B.C))/(Z*A)$.
                                                    NEGROOT(A.B.C)=(-B)/(2*A)-SQRT(DISC(A.B.C))/(2*A)$.
SORT
                     the
                              expression
                                                    DISC(A.B.C)=8.8-4.A.C$.
                                                    SQ(X)=((X)=(X))$.
              from
                      expanding the
                                                    SQPT(X)=((SQ((X)+4)+4*(X)1/(2*(4*((X)+4)))*
                                                                (((4*((X)+4))*(X))/(2*(SQ((X)+4)+4*(X)))):(S;
                                                  GLORAL EPROP. P1. P7:
                                                  ROUTINE POOT(X.Y.Z)=
sequence in R1
                        to the fourth
                                                    BEG1N
                                                      IF DISC(.X..Y..Z) LSS @ THEN EPROR+1 ELSE
                                                      IF DISC(.X .. Y .. Z) EQL @ THEN
term.
                                                        (R1--.Y/(Z+.X): R2--.Y/(Z+.X1)
                                                      ELSE (R1-POSPOOT(.X..Y..Z):R2-NEGPOOT(.X..Y..Z));
                                                    END:
                                                END ELUDOM
```

```
POSPOOT(A.B.C)=(-B)/(2+A)+SOPT(DISC(A.B.C))/(2+A)s.
                                                    NEGROOT(A.B.C)=(-B)/(2*A)-SQRT(DISC:A.B.C))/(2*A)$.
SORT is a macro identical to
                                                    DISC(A,B.C)=B+B-4+A+C$.
                                                    SQRT(X)=(XI+X; XJ+INFINITY:
                                                             WHILE (.XJ-.XI) GTP .EPSILON
the subroutine in R1.
                                                                DO (XI+.XJ: XJ+(.XI+.XI+(X))/(Z*.XI));
                                                              .xJ)s.
                                                    INFINITY-M777775
                                                  GLOBAL ERPOR.P1.R2.EPSILON:
                                                  ROUTINE POOT(X,Y,Z)=
                                                    BEGIN LOCAL XI.XJ:
                                                      IF DISC(.x..Y..Z) LSS @ THEN EPPOR-1 ELSE
                                                      IF DISC(.X..Y..Z) EQL 9 THEN
                                                        (R1+-.Y/(2+.X): P2+-.Y/(2+.X))
                                                      ELSE (R1-POSPOOT(.x..Y..Z):PZ-NEGFCOT(.X,.Y..Z)):
                                                    END:
                                                END ELUDON
```

BEGIN

EXAMPLES QUADRATIC FORMULA

The results of running KATE on R1, R2, and R3 are summarized in the following table:

	RI	R2	R3
NOOPT	86	196	108
ALLBUTSIM	52	42	62
ALLOPT	52	42	47

The output produced by KATE in the ALLOPT mode for each example is reproduced on the next three pages.

EXAMPLES OUADRATIC FORMULA

-- R1 --BEGIN MOCRO PDSPDDT(A,B,C)=(-B)/(2*A)*SQRT(DISC(A,B,C))/(2*A)\$. NEGROOT(A.R.C)=(-B)/(2*A)-SQRT(DISC(A.B.C))/(2*A)\$, DISC(A.B.C)=8+8-4*A*C\$: FORWARD SORT: GLOBAL ERROR R1 . R2: ROUTINE ROOT(X,Y,2)= BEGIN IF DISC(.X..Y..2) LSS 0 THEN ERROR+1 ELSE IF DISC(.X..Y..2) EQL 0 THEN (RI+-.Y/(2*.X); R2+-.Y/(2*.X)) ELSE (R1+POSRODT(.X.,Y.,Z):R2+NEGRODT(.X.,Y.,Z)); END: ROOT: .Y..Y..TS1 2 HUL 4 .. X .. T\$2 2 HUL .1\$2..2..1\$2 1 .T\$1..T\$2..T\$1 SUB 1 .T\$1.0..T\$2 LSS 9 T\$7.1\$1 BRE 1 LD FRROR.1.1 LSZ L\$1: 2..X..T\$2 DIV .Y..T\$2,-.T\$3 ž EQL BRE .T\$4.L\$3 R1..T\$3,-.T\$3 LDN R2..T\$3.-.T\$3 LDN LS4 DD 3 L\$3: PARM . T\$1 SQRT..TS0 CALL DIV .T\$0..T\$2..T\$0 SUB .T\$0..T\$3..T\$0 LD PARH R1..TS0..TS0 . T\$1 SQRT..TS0 CALL DPARK DIV .T\$0..T\$2..T\$0 .T\$3..T\$0.-.T\$3 ADD R2..T\$3,-.T\$3 12 L\$4: L\$2: 10 T\$0.0. 750 RTRN 1 ROUTINE SQRT(X)= LOCAL XI.XJ: GLOBAL EPSILON: MACRO INFINITY=#7777775; XI+.X: XJ-INFINITY: WHILE (.XJ-.XI) GTR .EPSILON DO (XI+.XJ; XJ+(.XI+.X)/(2*.XI)); .x.s END: SQRT:

2 4

LD

GTR

DDS

L\$6: SUB XJ.777777.777777

.XJ..XI..T\$1 .T\$1,.EPSILON,.T\$1

.T\$1.L\$7

```
LD
            XI..XJ..XJ
    HUL
            .XI..XI..TSI
    ADD
             .T$1..X..T$1
    MUL
            2..XI..T$2
             .TS1 .. TS2 . . TS1
    DIV
    LD
            XJ..T$1..T$1
    BP
            1 $6
Ĺ$7:
            T$0,.XJ,.T$0
    LD
    RTRN
                                         1
END ELUDOM
        TOTAL COST= 52
```

-- R2 --

```
BEGIN
   HACED
      POSPOOT(A.B.C)=(-B)/(2*A)+SQPT(DISC(A.B.C))/(2*A)$,
      NEGROOT(A.B.C)=(-B)/(2*A)-SQRT(DISC(A,B.C))/(2*A)$,
      DISC(A.B.C)=B+B-4+A+CS.
      50(X)=((X)=(X))$.
     SQRT(X)=((SQ((X)+4)+4+(X))/(2*(4*((X)+4)))+
                  (((4*((X)+4))*(X))/(2*(5Q((X)+4)+4*(X))))$;
   GLOBAL ERPOR.R1.R2:
ROUTINE PODT(X.Y.Z)=
     BEGIN
        IF DISC(.X,.Y,.Z) LSS 0 THEN ERROR-1 ELSE
        IF DISC(.X,.Y,.2) EQL 0 THEN

(R1+-.Y/(Z*.X); R2+-.Y/(Z*.X))
       ELSE (R1-POSPOOT(.X..Y..2);R2-NEGPOOT(.X..Y..2));
     END:
POOT:
    HUL
HUL
              .Y..Y..TS1
              4 . . X . . TS2
     HUL
              .152..2..152
     SUB
              .T$1..T$2..T$1
    LSS
              .151.0..152
                                               *
     BRE
              .752.L$1
    LD
             EPPOR-1-1
                                             ż
             L$2
                                                  3
1 $1 .
    HUL
             2..x..152
                                             2
    DIV
              .Y..T$2.~.T$3
                                            ž.
    EQL.
              .751.0..754
              .T$4.L$3
    LDN
             R1 .. T$3 .- . T$3
    LDN
             PZ..T$3.-.T$3
    BR
             LS4
                                                    3
1 $3
    ADD
              .751.4..754
.754..754..755
    HLI.
     HUL
              4..T$1..T$6
     ADD
              .185..186..185
              4.. 754, . 754
     HUL
              2..154..156
    DIV
              .155..156..156
     MUL
              .754..751..754
     MUL
              2..1$5..1$5
              .184..185..184
    DIV
                                             1
              .156,.154,.156
.156,.152,.156
     ADD
    DIV
     SUB
              .T$6..T$3..1$2
             R1..T$2..T$2
     ADD
              . T$3.. T$6.-. T$3
             R2..T$3.-.T$3
                                                    22
LS1:
L$2:
    LD
              TS0.0..TS0
    RTEN
 END ELUDOR
```

TOTAL COST= 42

```
-- R3 --
```

```
BEGIN
      MACED
        PDSP00T(A.B.C)=(-B)/(2*A)+S0PT(DISC(A.B.C))/(2*A)s.
        NEGRODT(A.B.C)=(-B)/(2*A)-SQRT(DISC(A.B.C))/(2*A)$,
        DISC(A.B.C)=B+B-4+A+CS.
        SQRT(X)=(X]-X: XJ-INFINITY;
               WHILE (.XJ-.XI) GTR .EPSILON
                   DD (XI=.XJ: XJ=(.XI=.XI)*(X))/(Z=.XI));
                . XJ)$.
        INF INITY=#777775:
      GLOBAL ERPOP.R1.R2.EPSILON:
     ROUTINE RODT(X.Y.Z)=
       BEGIN LOCAL XI.XJ
          IF DISC(.x..Y..Z) LSS 0 THEN ERPOR-1 FI SE
          IF DISC(.X..Y..Z) EQL 0 THEN
          (R1+-.Y/(2+.X); R2+-.Y/(2*.X))
         ELSE (R1-POSROOT(.x,.Y..2);R2-NEGROOT(.x,.Y,.2));
       END:
  ROOT:
      man
               .Y..Y..T$1
      MIN
              4 .. X . . TS2
      HUL
               .T$2..Z..T$2
      SUB
              .T$1..T$2..T$1
      LSS
              .T$1.0..T$2
      BRE
              .TS2.LS3
      LD
              ERPOR.1.1
      BR
              L$4
 L$3:
     HUL
              2..x..1$2
     DIU
              .Y..T$2.-.T$3
     EQL
              .T$1.0..T$4
     BRE
              .T$4.L$5
     LDN
              R1..753.-.153
     LDN
             R2..T$3.-.T$3
             L$6
                                          1
 L$5:
     SSCAL
             S$1.E$1
                                          1
 S$1:
     LD
             X1..TS1..TS1
     LD
             XJ.77777,77777
                                          2
1 69.
     SUB
             .XJ..XI..T$4
    GTR
             .TS4..EPSILON..TS4
     BRE
             .T$4.L$10
     LD
             X1..XJ..XJ
.X1..XI..TS4
     MUL
             .TS4..TS1..TS4
     noo
    MUL
             2 .. X1 . . T$5
     DIV
             .T$4..T$5..T$4
    LD
             XJ..T$4,.T$4
L$7
    BR
                                         1
                                                  14
L$10:
    DIV
             . XJ. . T$2 . . T$4
    RTRN
E$1:
    SUB
             .754..783..755
    LD
             R1..155..155
    SSCAL
            5$1..+1..7$4
    ADD
             .753..T$4,-.T$3
    LON
            R2..T$3.-.T$3
                                                27
L$6:
            T$0.0..T$0
END ELUDOM
```

TOTAL COST= 47

Notice that the S³ optimization had no effect on either R1 or R2. In the case of R3, on the other hand, a 25% improvement was realized by applying S³ optimization. The most interesting comparison, however, is between R1 and R3.

Both programs R1 and R3 represent the same logical structure to the programmer. The decision to declare SQRT as a macro or a routine does not effect that structure. Typically one expects the choice between the two is made in terms of some superficial estimate of the resulting time/space trade-off. The S3 optimization makes that same decision but more precisely. Indeed the S3 optimization did more than simply decide to open or close the SQRT computation in R3. The 10% reduction realized in R3 as compared with R1 results from:

- (1) not requiring parameters for S_1 since DISC(X,Y,Z) is available in T_1 and 2*X is available in T_2 , and
- (2) creating a strong similarity subroutine (S₁) for SQRT/DISC(X,Y,Z))/(2*X). Notice that this expression has no "logical identity" (as subroutine or macro) in the algorithm but S³, analyzing only the form of the program, identified it as a computational unit.
- Item (2) is the critical point. The results in this example and the examples which follow demonstrate that computationally coherent expressions (candidates for S3 optimization) do not necessarily correspond to the logically coherent expressions identified by the programmer as a macro or

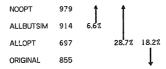
subroutine. Most discussion on optimization strategies which consider opening or closing subroutines has centered on examing those expressions which the programmer has identified as logically coherent. Similarity operates independently of the programmer's selection.

GADD-SUB

The second example comes from the Bliss-10 compiler. The routine GADD-SUB (abbreviated: GAS) generates code for add and subtract operations. The source and output from KATE's compilation of GAS in ALLOPT mode is reproduced in appendix A.

This version of GAS differs from the original version in the Bliss-10 compiler in that several of the macro declarations here were routines in the original. In particular, LITV, REGAK, TVRP, and RLITP were routines in the original version. The results of compiling GAS with NOOPT, ALLBUTSIM, and ALLOPT modes and of compiling the original with ALLBUTSIM are summarized as follows:

GAS



Again the difference in code size that results when S3 decides which expressions to close is striking. The table that follows is keyed to pages in appendix A and serves as a guide to locating the S3 optimizations in the output.

SSNAME	<u>SEMANTICS</u>	CALLS	COST	<u>PAGE</u>
\$1 \$2 \$3 \$5 \$6 \$7 \$10	GNEG(Y) LITV(n) RLIFP(n) RLEX(X) GANL(n ₁ ,NAMELEX(x),n ₂) GLTR(X) GASCOMMUTE	2 3 4 4 2 5	3 7 5 2 8 3	122 122 122 123 123 124 124
\$12 \$15 \$26 \$30 \$32 \$33 \$35	(X-GLTR(X);REGAK(X))123 REGAK(X) LITV(SLEX(n)) neq Ø TVRP(n) SIGN(X) GNEG(GAS(.ABSX,.ABSY,.ADDPOSSIBLE))	2 3 4 3 5 2 2	5 12 11 15 2 8 9	125 125 127 127 128 131 131

There are several observations to make about the results of S3. In the original source for GAS the routine REGAK was a single parmeter subroutine. The S3 optimization created a zero parameter subroutine S1s since all calls within GAS to REGAK passed the same parameter X. S6 is a case where S3 recognized that two calls on GANL passed the same second parameter and so created a new two parameter subroutine. S12 and S35 are examples of formally identical expressions which were not assigned a logical name (via macro or routine declaration) in the original source.

It is interesting to observe that the subroutines of the original text

were re-recognized as subprograms by \$3. One might ask why a good programmer would not have identified himself all the choices made by \$3. In the case of GASCOMMUTE, it would seem natural for the programmer to have made that identification. However, it is extremely unlikely that the same programmer would have identified \$6, \$12, \$33, and \$35 as code sequences to be closed although closing them did reduce code size by slightly less than \$42. More importantly, this example demonstrates that he need not be forced to make the choice between open and closed subprogram. An \$3 optimization can be used to perform this analysis.

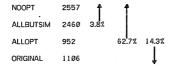
CPOLY

The final example is selected from the algorithms section of the Communications of the ACM[JT72]. CPOLY is a Fortran program to find all the zeros of a complex polynomial. Being a translation of an Algol procedure, it conformed easily to Bliss control syntax. In addition, the translation to Fortran had precluded recursive calls among the various subroutines. The source for CPOLY is reproduced in Appendix B.

CPOLY was transcribed into two Bliss versions with the body of SUBROUTINE CPOLY as the main body of the Bliss program. In one version the remaining subroutines were declared as macros. KATE compiled this program in NOOPT, ALLBUTSIM, and ALLOPT modes. KATE also compiled a second version in which the original subroutines remained as routines. The results are

summarized below:

CPOLY



The results are quite similar to those obtained in the GAS example. The large variation (62.7% vs. 3.8%) between invoking the S³ optimization and not invoking it results from the size of the subroutines involved and the frequency of the calls on them. The savings of the ALLOPT compilation over the ORIGINAL results primarily from two characteristics of the program:

- (1) Several of the subroutines, viz. SCALE, CAUCHY, NOSHFT, and FXSHFT were called only once. S³ simply compiled them in line.
- (2) Many of the procedures are passed parameters which are identical at all call sites. \$3 reduced calling overhead by removing those parameters

S3 OPTIMIZATION AND EXECUTION TIME

The preceding discussion on the effects of the S³ optimization has concentrated on reductions in code size. Since S³ reduced program size by

introducing subroutines, it is natural to assume that program execution time has been increased. In this section we will report on some preliminary analysis which demonstrates that such an assumption is not valid. We chose CPOLY for our analysis over GADD-SUB since the latter program is simply a large decision tree and has no loop expressions.

Our main difficulty in analyzing the effect of S3 on execution time is selecting a reasonable method for performing a static evaluation. For example, consider the problem of estimating the execution time of a branching control expression. There are three obvious alternatives: (1) select a particular branch, (2) average the execution times of the branches (assuming equi-probable selection of a branch), or (3) compute a weighted average of the branches (assigning a probability of selection to each branch). The added constraint that we intended to collect the data by hand compelled us to choose the first alternative and to limit our investigation to the inner loop of CPOLY which we identified as the loop in VRSHFT called from FXSHFT.

Two control paths through VRSHFT and the subroutine called by VRSHFT were selected. The first path was chosen by selecting those branches which entail the largest number of instructions. That is the longest (deepest) path through VRSHFT and its calls. The number of instructions executed in the original version was 3630±NN + 10360 and for the S³-produced version 3810±NN + 6090. The parameter NN is the degree of the input polynomial plus one and is constrained to be ≤50. Thus in the worst case (NN=50) the

 S^3 version requires 2.5% more execution time than the original. As NN decreases the performance of the S^3 version improves. If NN=10, then the S^3 version requires 5.5% less execution time than the original.

The second control path which we selected was shorter (i.e. fewer instructions per iteration): The original version executed 860°NN + 3030 instructions whereas the S³-produced version executed 860°NN + 1840. The NN-terms in the equations are identical since no S³-created (and not specified by the programmer) subroutines were executed in the NN-dependent loops. If NN=50, then S³ reduced execution time by 2.5%; and if N=10, then S³ reduced execution time by 10%.

The effect of S³ optimization on the execution time of a program clearly requires more study than that given by this preliminary analysis. The purpose of presenting the results of this initial investigation is to dispel the assumption that S³ optimization necessarily increases the execution time of a program. Indeed that had been our assumption before we studied the effects of S³ on CPOLY more closely.

SUMMARY

Having produced a set of numbers measuring the effects of the program KATE on a few examples, it is important to place this information in the proper perspective. Chapter II introduced a collection of primitives used

to describe the class of classical optimization techniques. The effectiveness of those optimizations is not an issue to this thesis. The success of the Fortran-H experiment which embodies those optimizations has already verified their utility. The merit of Chapter II lies in the concise statement of these optimization strategies and a correspondingly simple implementation of them.

The similarity notion, on the other hand, is a new concept. Chapter III described a number of optimizations in terms of similarity. We selected one of those, the S3 optimization, and implemented it in KATE. S3 was selected because it dealt with an area of object code optimization not touched by Chapter II — the opening and closing of subprograms. Cocke and Schwartz discuss this area in some detail. However they concentrate on working with subprograms already identified by the programmer rather than on discovering the subprograms independent of the programmer. In addition, they only consider opening subprograms and reducing the amount of linkage code. The results that KATE produced are not to be interpreted as conclusive evidence that. S3 optimization will produce a 10% to 15% reduction in program size across the board. The results do say that S3, which is concisely and coherently describable in terms of the similarity notion, has potential for producing significant reductions in object code size.

Finally, if one examines any of the above examples, he can find places where KATE could have done better or where, if the original program were

restructured, \$3 would not have produced the same favorable results. We do not propose a contest between programmer and compiler to discover some "minimal" program. We see the \$3 optimization in the following light. Let the programmer design the logical structure of his program and identify his computational sequences on the basis of their logical coherence. An \$3 optimization can decide for him between implementing those sequences as closed or open subprograms.

CHAPTER V CONCLUSION

This final chapter is divided into two sections. The first section summarizes the results of our investigation. The second part suggests future directions in which this study can progress.

SUMMARY OF THESIS RESULTS

Chapter II motivated, defined, and used a collection of concepts for describing code motion, redundant expression elimination, and strength reduction optimizations. The concision of those decriptions demonstrates that the goal of discovering a set of primitives sufficiently powerful to enable concise descriptions of a class of optimizations has been achieved. Furthermore, the descriptions are independent of the intermediate representation of the program. Language independence has been accomplished by isolating language-dependent characteristics in the ordering relations (a, \prec, \ll, \prec) . Finally, although the optimizations themselves may on the surface appear to be unrelated, the primitives provide a homogeneous description which, in turn, leads to a compact, cleanly structured implementation.

A new concept, similarity, was introduced in Chapter III. A collection of new optimizations was defined in terms of the similarity notion. One of these new optimizations, S3, was examined in greater detail. The discussion in Chapter III (and the analysis in Chapter IV) demonstrates that S3 opens a significant new area of investigation into program optimizations. Previous research in optimization has done very little in the area of optimizations involving subprograms. No work, known to us, has investigated the possibility of using a compiler to determine the computational units to be implemented as closed subroutines.

FUTURE RESEARCH

In the process of doing this research a number of areas of possibile future investigation have emerged. Some of them are short-range and reasonably well-defined while others are long-range and less specific.

The program KATE implemented the primitives of Chapter II and the similarity function, SIGMA, defined in Chapter III. The evolution of the primitives and the construction of KATE proceeded in parallel during our investigation. Each process provided information for the development of the other. However the major emphasis lay in the development of the primitives. Now that the primitives have evolved to their present state, it would be worthwhile to reconstruct KATE and observe the effect on the resulting program. Since opsimizing compilers are noted for being

expensive in terms of both time and space, one might conentrate on examining alternate implementations of the primitives which reduce this overhead.

Chapter III developed a particular similarity function, SIGMA. That function was evolved with the S3 optimization technique in mind. It is not clear that SIGMA is the appropriate similarity function for all the optimizations defined in Chapter III. An obvious area of investigation lies in discovering other useful similarity functions. Particularly, one might examine similarity functions which are sensitive to execution time overhead and the use of temporary storage. The set of optimizations described in terms of the similarity concept in Chapter III are new. In addition to developing new similarity functions, there is certainly the potential for discovering more optimizations defined in terms of similarity.

Another area of investigation is related to the notion of strength reduction. In Chapter II we began the section on strength reduction by posing the problem of discovering a relation F, such that $F(e,\Delta e) = e^t$ and the cost of evaluating $F(e,\Delta e)$ is less than the cost of evaluating e^t . The statement of this problem is motivated by the observation that strength reduction seems too specialized. The restriction to polynomials and looping environments is reasonably restrictive. The thesis described the feasibility of strength reduction optimizations in non-looping environments. The generalization to non-polynomial expressions, on the

other hand, remains an open question. The problem consists of discovering a set of non-polynomial expression pairs (e,e') for which there exists a closed-form relation F satisfying the equation $F(e, \Delta e) = e'$,

Finally and, to our mind most importantly, a spectrum of questions opened by the S³ optimization technique remains to be studied. S³ was developed in the context of an investigation into object code optimization. Indeed one area of study is an investigation into modifications to the heuristics implemented in SiGMA and reconsideration of the overall structure of the S³ module in KATE. There are, however, other directions to be pursued.

At the end of Chapter IV we presented a brief summary of a preliminary investigation into the effect of S³ optimization on the execution time of a program. That investigation suggests two area for future study. First, there is the problem of performing a static analysis on the execution time of a program. Can one determine a meaningful data-independent measure of execution time? Can a program be analyzed to determine the kind of information that must be known about the input data in order to perfrom a valid analysis? Presumably a programmer makes some assumptions about the data input to a program in order to decide among alternative algorithms. Perhaps those assumptions can be incorporated into a static analysis of execution time. Second, the function SIGMA was designed to minimize object code size. How does one design a similarity function that is more sensitive to execution time? There are obvious parameters like loop depth

and calling overhead. It seems clear, however, that heuristics encoded in a execution-time-sensitive similarity function require the same kind of information used in a static evaluation of execution time. Hence these two areas appear to be closely related.

In analyzing the form of a program, KATE discovers a set of computationally coherent expressions. Our initial investigation into this area, discussed in Chapter IV, demonstrated deviations from the selections made by the programmer. It is interesting to consider what one might learn about the structure of programs by analyzing the results of allowing an S3 pass to select subprograms. Will S3 consistently outperform the programmer in terms of reducing program size? Do the (potentially) different subprograms selected by S3 provide significant feedback on the programmer's choice of logically coherent subprograms?

Some current research by S.L. Gerhart[GE72] involves the verification of APL programs. One aspect of this work is concerned with investigating the effect of the powerful APL operators on the verification process. For example, one observes that an algorithm represented by a nested-loop expression in Algol can perhaps be represented by a single operator in APL. As a result, a verification of the APL program should proceed with less difficulty than the verification of the corresponding Algol program since the effect of the involved Algol control expressions has been captured in a single operator. Intuitively this models a mathematician's approach to generating a large, involved proof. He typically identifies a set of

sub-goals (lemmas -- macros -- subroutines). Having verified the sub-goals, he proceeds to combine these into a verification of the original theorem. It seems promising, then, to investigate the usefulness of similarity for discovering sub-goals and thereby reduce the complexity of the verification process.

These last two suggestions are not directly related to the area of object code optimization, but are natural outgrowths from observing the effect of S3. They offer a wide range of interest for future study.

DECIM STRUCTUPE VECTOR(1)=(.VECTOR+.1); LEFTF=0s. RIGHTH=#1777775. NEGF=18. POSNSIZEF=26. ADDI=#271\$. SUBI=#278 LTEF=1\$. SUB1=#275\$. ARTEF=05. STACKUAR=1110+1155. ADDM=#273\$. SUBH=#276\$. ADD=#270\$. SU8=#2745. DOTM=11135. 1.5Me1175. NEGH=11158. PTFHum3275 LSSTEM=#37718\$, STEM=#377\$, RTESTEM=#1777775. ZERO=#4000\$. ZERO36=36\$. POSNS12EH=#1777775; SIGN(X)=((X) AND NEGHIS. LITP(X)=(((X) AND NOT STEM) EQL 0)s. NAMP(X)=(((X) AND NOT STEM) EQL (LSM OR ZERO36))\$. GABS(X)=((X) AND NOT NEGM)S. PERCNAMPIX)=(((X) AND NOT STEM) FOL ISHIS. STOCKHOPPINI-STOCKHOPHI- STINIE. LITU(X)=(IF (X)(UFF) THEN ATT(X)(LTFF)) FLSE (X)(LTFF))S. REGAK(X)=(IF NOT (.RT(:X)(RTEF)) AND (X)(DTF)) THEN .RT((X)(RTEF)) ELSE GMA(X))\$, CODE(F,A,M)=(INST-(F)127 OR (A)129 OR (M))s, TURP(X)=(IF REGP(X) THEN ITRP(X) OR (.RT((X)(RTEF1) EQL .OPTTOREGROOR) ELSE 015: GLOBAL OPTTOREGADOR.RT.INST.LT.ST: EXTERNAL SPOUTINESS PCIVE.GNEG.LITLEXEME.GANL.MPTRTYP.GLTR.GMA.TVMP.DCRP. GLAR.GLTH.PEGSEARCH.SHOULDEXCH.PEGAR.MEMORYA.GCLTR.REGP. READY . ITRP: ROUTINE GAS(X,Y,F)= IGENERATE CODE FOR XEY WHERE & IS CASE F OF SET +:- TES. THIS IS UNDOUBTEDLY THE BEST (MORST?) CASE FOR SHOWING THE ! COMPLEXITY OF THE DELAYING MECHANISM. IT HOULD BE FAIR TO SAY ! THAT THIS POUTINE IS BIASED TOWARDS OPTIMIZING STRUCTURE ACCESSING. ! I.E. FODITION BY INDEXING. FOR EXAMPLE, HAEN PASSED THE OPERANDS ! FOR .A * 1. GAS LOADS .A INTO A PEGISTER (SAY R) AND RETURNS A LEXEME ! OF THE FORM (.R+1) (I.E. RTEF=R AND LSSTEF=1). THE IDEA HERE IS THAT ! IF THE EXPRESSION .A + 1 HAS APPEARED IN THE CONTEXT *(.A+1)<9.36>+EXP* ! THEN THE ADDITION HOULD BE ACCOMPLISHED BY INDEXING IN THE INSTRUCTION: "MOVEH EXP.1(P)." THE SET OF SPECIAL CASES IS COMMENTED ON THE PIGHT SIDE OF ! THE CODE. E.G. !(@P+N)+L IS TO BE INTEPPRETED TO MEAN: X= LEXEME REP. REG + NAME Y= LITEPAL L F= +. FOLLOWING THE SET OF SPECIAL CASES THE ROUTINE ATTEMPTS TO ! HANDLE THE EIGHT CASES THAT ARISE FROM F AND THE POSSIBILITY OF ! UNARY MINUS ON X OR Y OR BOTH. (2) X-Y (1) X+Y (3) X+-Y (1) X--Y (5) -X+Y (6) -X--Y (7) -X+-Y (B) -X-Y MACPO GASCOMMUTE=(GAS(IF .F THEN GAEG(.Y) ELSE .Y..X AND NOT NEGH..X(NEGF)));

PHOTO DESCRIPTIVE (X MIG NOT PRESENT DO, 6 MIG (X MIG REEN) MED 0);

HOUSE PLEATING (X MIG NOT PRESENT)

HOUSE (X MIG NOT PLEATING (X MIG NOT PLEATING X PROPERTY IN THE NOT PLEATING X PROPERTY PROPERTY IN THE NOT PLANTY PROPERTY PR

ABSY, ! GABS(.Y)
ABSX, ! GABS(.X)
XPEG.YREG.R.

```
! F EQL SIGN(-Y)
  ADDPOSSIBLE:
HACRO
   TEMPX=R[0]S.
                  ! X IS A TEMP REG
   TEMPY-RIIIS: ! Y IS A TEMP PEG
MAP VECTOR X.Y.
PCIVR(.x,.Y);
ABSY+GABS(.Y); ABSX-GABS(.X);
IF LITP(.Y) THEN
                                                                          !X-L
  IF .F THEN GAS(.X,GNEG(.Y),0) ELSE
                                                                          1440
  IF .Y EQL ZERO THEN .X ELSE
  IF LITP(.X) THEN
                                                                          !L+L
     LITLEXEME(LITV(.X)+LITV(.Y)) ELSE
  IF PLITP(.ABSX) THEN
                                                                          !(mR+L)+L
    GAS(SLEX(.X)..Y..X(NEGF)) OR (.X AND (NEGH OR RTEH)) ELSE
  IF NAMP(.X) THEN
                                                                          !N+L
     GANL(8..X..Y) ELSE
TE PNAMP(.X) THEN
                                                                          !(eR+N)+L
     GANL (PLEX(.X).NAMELEX(.X)..Y) ELSE
                                                                          !X+L
  IF (IF ZEPONAMP(.X) THEN
       BEGIN
         YVALUE+LITY(.Y):
         (.YVALUE AND RIGHTM) EQL 8
           AND NOT STACKVARP( .X(STEF1)
       FND
     ELSE 0) THEN
                                                                          IXCR. 8>+L
    MPTRTYP(.YVALUE(LEFTF)..X) ELSE
  GLTR(.X) OR .Y
  ELSE
IF LITP(.X) THEN
                                                                           IL+Y
   GASCOMMUTE ELSE
IF ZERONAMP(.Y) THEN
                                                                           TXEY(B.B)
    (CODE:CASE .F OF SET ADDI:SUBI TES, (X-GLTR(.X):REGAK(.X)), GMA(.Y OR DOTH));
    .X1 ELSE
IF ZERONAMP(.X) THEN GASCOMMUTE ELSE
REGIN
                                                                           !x<8.8>£Y
ADDPOSSIBLE+.F EQL SIGN(.Y);
IF NAMP(.ABSY) AND .ADDPOSSIBLE THEN
  IF REGP(.X) THEN
                                                                           !eR+N
      .X OR (.ABSY AND LISTEM) ELSE
  IF RLITP(.X) THEN
                                                                           LIADAL 14N
      GANL (RLEX(.X), .ABSY, SLEX(.X)) ELSE
                                                                           1X+N
  GLTP(.X) DR (.ABSY AND LSSTEM)
  ELSE
 IF NAMP(.ABSX) THEN
                                                                           !NEY
    GASCONMUTE ELSE
 IF PNAMP(.ABSX) THEN
  IF (IF PLITP(.ABSY) THEN LITV(SLEX(.Y)) NEQ 0) AND .ADDPOSSIBLE THEN
                                                                           !(eR+N)+(eR'+L)
      REGIN
      IF TURP(PLEX(.X)) THEN
        (XPEG+PLEX(.X);YPEG-PLEX(.Y))
      ELSE (XPEG-PLEX(.Y):YPEG-PLEX(.X));
      GAS(GANL(.XREG.NAMELEX(.X).SLEX(.Y))..YREG.0) OR (.X AND NEGM)
      END
    ELSE
                                                                           !(#R+N)&Y
  GAS(GAS(PLEX(.X)..Y..F).(.X AND NOT RTEM) OR ZERO36.0)
```

END ELUDON

```
ELSE
   IF PNAMP(.ARSY) THEN
                                                                                 IXE(aP+N)
     GASCOMMUTE ELSE
   IF (IF PLITP(.ABSX) THEN LITU(SLEX(.ABSX)) NEQ 0) THEN
                                                                                 ! (#R+L)&Y
     BEGIN MACRO X1=ABSXS:
      X1-GAS(.X AND NOT LSSTEM .. Y .. F):
      IF .XINEGF) AND .XI(NEGF) THEN
        GNEG(GAS(SLEX(.X).GABS(.X1).0)) ELSE
      GAS(IF .XINEGF) THEN GNEG(SLEX(.X)) ELSE SLEX(.X).GABS(.X1)..XI(NEGF))
     END ELSE
   IF (IF RLITP(.ABSY) THEN LITV(SLEX(.ABSY)) NEQ 0) THEN
                                                                                 IXECAPAL S
     GASCOMMUTE ELSE
   IF TUMP(.Y) AND DERP(.X) THEN
     (CODE(IF .ADDPOSSIBLE THEN ADDM ELSE SUBM.
           (X-GLAR(.X):PEGAK(.X)).GMA(Y-GLTM(.ABSY))); .Y) ELSE
   IF TUMP(,X) THEN
                                                                                 !MS.Y
     GASCOMMUTE ELSE
   BEGIN
   REGSEARCH(X.Y):
   ABSX-GABS(.X): ABSY-GABS(.Y);
   IF (TEMPX+TURP(.ABSX)) AND (TEMPY+TURP(.ABSY)) THEN
      IF SHOULDEXCH(.X,.Y) THEN
        GASCOMMUTE ELSE
       IF SIGN(.X) THEN
         IF .ADDPOSSIBLE AND .RT(.X(RTEF))(ARTEF) NEQ .VREG THEN
                                                                                 15.6
           GASCOMMUTE ELSE
                                                                                 17.8
         GNEG(GAS(.ABSX,.ABSY,.ADDPOSSIBLE))
       (CODE(IF .ADDPOSSIBLE THEN ADD ELSE SUB-
              (X+GLTR(.X):REGAK(.X)).REGAR(GLTR(.ABSY)));
                                                                                11-4
       .x)
    ELSE
  IF . TEMPX THEN
     IF SIGN(.X) THEN
                                                                                15-8
       GNEG(GAS(.ABSX..ABSY..ADDPDSSIRLE)) FLSE
     (CODE(IF .AODPOSSIBLE THEN ADD ELSE SUB,
           (X+GLTR(.X);REGAK(.X)),
                                                                                11-4
           MEMORYA(.Y));
    . x)
  ELSE
IF . TEMPY THEN
     GASCOMMUTE ELSE
  IF SIGN(.X) THEN
    IF .ADDPOSSIBLE THEN
                                                                                15-6
        GASCOMMUTE ELSE
                                                                                17-R
     BEGIN
     X+GOLTR(.X);
     IF SIGN(.X) THEN
      GNEG(GAS(GABS(.X)..ABSY.0)) ELSE
    GAS(.X..ABSY.1)
    END FLSE
                                                                                11-6
  IF PERDY(.X) THEN
    IF .ADDPOSSIBLE THEN GAS(GLTR(.ABSY)..X.0) ELSE
IF PEADY(.ABSY) THEN GAS(GLTR(.X),.ABSY,1) ELSE
    GNEG (GAS (GLTR( . ABSY) . . X.1))
    ELSE
  GAS(GLTR(.X),.ABSY..F OR SIGN(.Y))
  END
  END
END:
```

GAS:						
PAF	ж .х		1			
PAF			i		•	
CAL	L PCIVE.	.T\$0	1	*		
DPF			1	. •		
ANE		0001T\$1	2			
LD		IS1TS1	1	*		
ANI		0001T\$1	. 2	•		
LD		151151	1 2	:		
EQL	. 1\$3.0		é	-		
BRE			ĭ			
BRE			3			
PAR			1			
550	AL 551.ES1	ı	1	=		
551:						
PAR			1	-		
DPF		150	1	:		
PTE			i	:		
E\$1:			•	-		
PAR	M .TS9		1			
PAR	m 0		1			
CAL		60	1			
DPF			. 1	*		
, LD	T\$4,.T\$	10,.154	1			
BR	' LS4		. 1	12.		
L\$3:				12.		
EQL	. Y.4000)T\$6	2			
BRF		5	ī			
LD	7\$5,.X.	. T\$5	1			
BR	L\$6		1			
				2		
L\$5:			_			
AND EQL			- 2	:		
BRE	.T\$10.L		1	- 2		
LD	T\$10X	T\$10	i	_		
SSC			i			
5\$2:						
ADD			2		*	
ADD					×	
BRF			1		٠.	
ADD LD		2T\$13 T\$13,.T\$11	1		:	
BR	L\$12	1913,.1911	i			
OK.			•		3	
L\$11:						
LD	T\$11T	\$12T\$11	1		•	
					1	
L\$12:						
. RTF	M				•	
LD LD	T\$10Y	TE18	1			
550			î			
ADD	.T\$11	T\$11T\$12	2			
PAR			1			
CAL		ME T\$0	-1		=	
DPF			1		• .	
LD BR	TS6TS LS10	10156	1		:	
(BR	1510		1		19	
LS7:					.5	
LD	T\$13A	85×T\$13	. 1		. `	
SSC			i			
583:						
- AND		200000T\$14	2		*	
EQL			8		* .	
AND NEG		771\$15	2		:	

RIEN		1	•
E\$3: BPF	.T\$14.L\$13	2	
eND	.X.177777T\$15	ž	
PAPH	.T\$15	ī	
PAPH	.Y	1	`*
ADD	X.1T\$15	0	*
PAPM	7\$15	1	•
CALL DPARM	GASTSO	1	
END	.x.100377T\$15	2	
OR.	.T\$0T\$15T\$0	ì	
LD	T\$12T\$0T\$12	1	
BR	L\$14	1	*
L\$13:			12
EQL.	.T\$7,244T\$16	1	
BRF	.T\$16.L\$15	i	
PARM	0	i	
PAPM	.x	1	•
PARM	.Y	1	
CALL	GANLTS0	1	:
DPARM LD	3 7\$157\$07\$15	1	:
BR	L\$16	i	- :
	1310	•	7
L\$15:			
ADD	X.2T\$20	9	=
EQL	T\$20.0T\$20	1	
SSCAL	.1520.L521 555.E55	1	٠.
SSCRL SSS:	222-522	1	•
AND	.x.377,.T\$Z1	z	
RIPN		ĭ	
E\$5:			
NEO	.TS21.0TS2Z	8	*
BRF	.T\$22.L\$23	1	
DR OND	.T\$7.44,.T\$22 .T\$22,-400,.T\$22	. Z	
EQL	.1922.244,.1922	i	
LD	1\$201\$221\$20	i	
BR	L\$24	i	
L\$23: LD	T\$20,0,.T\$20	. 1	
LD	1520.01520		
L\$24:			
LD	T\$17T\$20T\$17	1.	
BP.	LSZZ	1	
			111
L\$21: LD	T\$17.0T\$17	1 .	
LU	1917.05.1917		ī
L\$22:			
BPF	.TS17.LS17	1	
SSCAL	5\$5+17\$21	1	
LD	T\$22T\$21T\$22	1	
LD	T\$23YT\$23	1	*
SSCAL SSG:	5\$6.E\$6	1	
PARM	.1522	1	
AND	.X.177400T\$24	ž ·	
OP.	.1524.441524	ī	
PAPM	.T\$24	1	
PAPM	.1\$23	1	*
CALL	GANLTSO	1	
DPAPH RTRN	3	1	:
E\$6:			. •
LD	7515T\$0T\$16	1	
. BB	L\$20	i	*
			19
L\$17:	707 700 707		

LD SSCAL LD AMD EOL ADD ADD RISH ADD ANDCR LD BR	1510Y1510 552+11511 YVALUE15111511 .YVALUE.1777771511 .T511.0151 X.3157 51157 51157 .T51157 .T51157 .T51157 .T51157 .T51157 .T51157 .T51157 .T51157 .T51157	1 1 2 0 0 1 3 1 1	. *
L\$27: LD	T\$26.0T\$26	1	
L\$30: BPF ADD PARH PARH CALL DPARH LD BR	.1526.L925 YVALUE.O1526 7526 .X HPIRTYP,.TS0 2 1525TS0,.TS25 L526	1 9 1 1 1 1	i
L\$25: SSCAL	S\$7.E\$7	1	
PARM CALL DPARM RTRN	CLTRTSG	1 1 1 1	
E\$?: OR LD	.150Y150 T\$25T\$0T\$25	1	
L\$26: LD	T\$16T\$25T\$16	1	
L\$20: LD	T\$15T\$16T\$15	1 .	65
L\$16: LD	T\$12T\$15T\$12	1	# 75
LS14: LD	T\$6T\$12T\$6	1 .	98
L\$10: LD L\$6:	1\$51\$6,.1\$5	1	121
LD LS4:	T\$4T\$5T\$4	1	127
LD BR	T\$2T\$4T\$Z L\$2	1	144
LS1: AND EQL BRF SSCAL SS10:	.x400155 .155.0156 .156.L531 S\$10.E\$10	2 8 1 1	:
BRF SSCAL LD BR	.F.L\$33 5\$1*1T\$0 T\$6T\$0T\$5 L\$34	3 1 1 1	* * *
L\$33: LD	, T\$6Y,.T\$6	1	1
L\$31: PAPM	.1\$6	1	•

APPENDIX A

	***	1	
PAPIT	.TS1 X.1TSG	é	:
PARM	1\$6	1	•
CALL	GAS1\$0	1	:
RIPN	3	1	
E\$10:			
LD BR	T\$4T\$0T\$4 L\$3Z	1	:
			16
L\$31: EQL	.1\$3.2001\$3	1	
BRF	.1\$3.L\$35	í	
XCT	.F.L\$37	ž	
L\$37:	' 618	1	_
BR	L\$41	i	
L\$10:		-	
LD	T\$3.271T\$3	1	
BR	L\$12	1	2
L\$41:			
LD	T\$3.275T\$3	1	1
L\$12:			. 1
LTSH	.753.33753	1	
SSCAL	S\$12.E\$12	1	
5\$12:			_
SSCAL LD	5\$7*1T\$0 X,.T\$0T\$0	1 1	:
SSCAL	5\$15.E\$15	i	
S\$15:			
ADD	.K,4,.T\$15 RT,.T\$15T\$15	1 0	:
ADD ADD	.X.5T\$16	ž	:
AND	7\$157\$167\$16	1	
BRT	.T\$16,L\$43	1	
LD	T\$12,T\$15,T\$12	i 1	
BR	L511	1	2
L\$43:		_	
PARM	.X GMAT\$0	1	
DPARM	1	i	
LD	T\$12T\$9T\$12	1	:
L\$11:			4
RTEN		1	
E\$15:		_	
LTSH	.T\$12.27T\$16	2	:
E\$1Z:		•	
OP:	.1\$3,.1\$161\$3	1	•
OP.	.Y.20000T\$25	2	•
PARM CALL	.T\$25 GMAT\$0	1	
DPARM	1	i	:
OR	.153150153	1	
LD	INST1\$31\$3	1	
ED BR	T\$6XT\$6 L\$36	1	:
Bic	C+35		37
L\$35:			_
- BRF	.T\$5,200T\$25 .T\$25,L\$45	1	:
SSCAL	S\$10+1T\$0	i	٠.
LD	T\$3T\$0T\$3	i	
BR	L\$46	i	
L\$45:			3
AND	.Y.100000T\$25	2	
EQL	.F16251825	1	
LD	ANDPOSSIBL T\$25 T\$25	1 2	:
AND EQL	.ABSY,-400T\$26 .T\$26.244T\$11	1	
L.W.L		•	

	AND	.T\$11ADDPOSSIBLT\$11	1	•
	PAPH	.T\$11.L\$47	1	•
	CALL	REGPTS0	i	
	DPAPH	1	i	:
	BPF	.TS0.LS51	ż	
	END	.ABSY.177400T\$7	2	
	OP.	.X1\$77\$?	1	
	LD	T\$11T\$7T\$11	1	
	BR	L\$52	1	•
				5
	51: AND	.x2000001\$27 ·	2	
	EQL	.1827.01827	6	i i
	SSEAL	585417\$21	ĭ	
	NEO	.1521.01530	ē	*
	AND	.1\$271\$301\$27	i	
	BRF	.T\$27.L\$53	1	
	PARM	. 1521	1	
	PARM	.ABSY	1	
	AND	.X.177777,.T\$27	2	
	PAPM	.1\$27	1	
	DPARM	GANLTSO 3	1	*
	LD	187180187	1	:
	BR	L\$54	· .	
	Dr.	2001	,	9
	53:			
	SSCAL	557,.+1,.750	1	
	AND	.AB5Y.177400T\$27	ż	
	0R	.1501527150	1	
	LD	187180187	1	
				. 5
LS	54:	T\$11T\$7T\$11		_
	LD	191119771911	1	20
Ls	52:			
	LD	T\$25T\$11T\$25	1	
	BR	L\$50	1	
				32
LS	47: AND	.ABSX,-108T\$?	2	
	EQL	.T\$7.244T\$27	í	:
	BRF	.T\$27,L\$55	i	- 1
	SSCAL	5\$10+1750	i	٠,
	LD	T\$11T\$0T\$11	i	
	BR	L\$56	i	
				3
. 5	55:			
	ADD	AB5X.2T\$31	9	
	EOL	T\$31.0 T\$31	1	
	END BISI	. T%31 -L\$61	1 2	•
	NEO	.1\$13.3771\$32 .1\$32.0,.1\$32	9	:
	BRE	.T932.L\$63	ž	- :
	DR	.157,11,.157	ī	_
	AND	.T\$7,-400T\$7	i	
	EQL	.787.244787	1	
	LD	T\$31T\$7T\$31	1	
	BR	LS64	1	
. 5	63: LD	T\$31.0T\$31	1	
	LU	1831.01831	1	
•	64:			
•	LD	T\$30T\$31T\$30	,	
	BR	L\$62	i	
į			-	17
Š	61:			
	FD	T\$30.0T\$30	1	
				. 1
. 5	62: BPF	.7830.L\$57	1	-
	SSCAL	555+11521	;	٠.
	LD	1\$13ABSYT\$13	i	:

551			1	•
BRI LD	157Y		2	٠,
	AL 5526.ES2		î	
5\$26:			_	
ANI ADI		777T\$34 .T\$35	2	` :
AN	.157.177	7771\$34	2	
ADI			9	:
BRI		71	1	•
LD	T\$33T	\$361\$33	i	
BP.	L\$72		1	
L\$71:				
LD.	T\$33T\$	351\$33	1	
L\$72:	.0.5533.0	Test	0	
RT			ĭ	
E\$26:				
LD BP:	T\$31T\$ L\$70	357\$31	1	
	23.0		•	16
L\$67:			1	_
LD	T\$31.0	1\$31	1	-
L\$70:				
AN	.T\$31A	DDPOSSIBLT\$31	1 .	
BR eN			1 2	
LD			1	
	CAL S\$30.E\$3	9	1	*
S\$30: PA	PH .T\$37		1	
CA	LL REGPTS	0	1	
DP	ARM 1		1	
BR PA		5	2	
CA		0	i	
DP	ARM 1	•	1	
AD		.1541	1	
AD EQ		0P110PEGAD T\$41	0 2	
OP.	.T\$0T\$	41T\$0	1	
LD		0T\$40	1	
BP	L\$76		1	
L\$75:				
LD	T\$10.0	T\$40	1	
L\$76:				
RT	RN		1	
E\$30:		-72	2	
BR LD	F .T\$40.L\$	21 T\$21	1	-
LD	YPEGTS	317\$31	1	
BR	L\$74		1	
L\$73:				
LD		317\$31	1	
LC	YREG TS	217\$21	1	
L\$74:				
LD		EG1\$22	1,	
AN		7T\$31 i31T\$23	2	:
L0 59	CAL 5\$6+1.	.750	ì	
PF	RM .150		1 .	
	RM .YREG		!	:
	PM 0 LL GASTSG		1	:
DF	ARM 3		1	
59	CAL S\$32.E\$3	32	1	

			_		
	AND RTRN	.X.100000T\$31	2		- :
E\$3	2:				
	DR LD	.T\$0T\$31T\$0 T\$30T\$0,.T\$30	1		:
	BP.	L\$66	i		
	_			•	11
L\$6	PARM	.T\$21	1		
	PARH	. Y	i		
	PARM	.F	. !		
	DPARM	GAST\$0 3	1		
	PARM	. T\$0	1		
	DR PARM	.T\$5.44T\$5 .	1		*
	PARM	0	i		÷
	DPARM	GAS. TSO	1		
	LD	3 T\$30T\$0T\$30	1 1		:
		140011140111400	•		12
L\$6	6: LD	T\$27T\$30T\$27	,		*
	BR ED	L\$60	;		÷
			-		84
L\$5	7ı ADD	ABSY.2T\$41	8		
	EQL	7541.0 7541	î		:
	BRF	.T\$41.L\$101	1		
	AND NEO	.ABSY.377T\$42	2		
	BRF	.T\$42.L\$103	z		
	OR AND	.T\$26.44T\$26	1		
	EQL	.1\$26,-400,.1\$26 .1\$26,244,.1\$26	1 1		
	LD	7\$41,.7\$26,.7\$41	1		
	BR	L\$104	1		
L\$1	03:				
	LD	T\$41,0,.T\$41	1		
LSI	04:				
	LD	T\$5T\$41,.T\$5	1		-
	BR	L\$102	1		12
LSI					
	LD	T\$5.0T\$5	1		1
LSI	32:				1
	BRF	.TSS.LS77	1		
	SSCAL	\$\$10+1,.1\$0 7\$307\$0,.7\$30	1		:
	BR	L\$100	i	2	
L\$7			4		3
		T\$13ABSXT\$13	1		
	SSCAL	S\$3+1T\$14	1		
	BRF LD	.T\$14.L\$107 T\$7ABSXT\$7	2		٠.
	SSCAL	S\$28+11\$35	:		
	LD	T\$41,.1\$35,.1\$41	1		
	BR	L\$116	1		
L\$16					
	LD	T\$41.CT\$41	1		1
L\$11	10-				,
	3PF	.T\$41.L\$195	1		
	AND FAPH	.X177401T\$41 .T\$41	2 1		:
	PARM	.Y	i		
	PARM	.F	1		•
	DPAR:	GAST\$0 3	1		:
	LD	ABSXT\$0T\$0	i		

			The state of the s
AND	.X.177777T\$35	2	•
AND	.ABSX100001T\$26	2	*
ADD	X.1T\$42	е.	•
ADD	ABSX.1T\$43	Θ.	
AND	1\$421\$43 1\$42	2	•
BRF	.T\$42,L\$111	1	
PARH	. T\$35	1	
PARH	.1\$26	1	•
PARM CALL	0	1	• .
DPARM	GAST\$0 3	1	
PARM	.750 .	1	
CALL	GNEGTSO	1	•
DPARM	1	i	
LD	T\$41T\$0T\$41	1	
BR	L\$112	i	
U.		•	10
L\$111:			
BRF	T\$42,L\$113	3	
PARM	.T\$35	í	
CALL	GNEGTSO	i	
DPARM	1	i	
LD	T\$42T\$0T\$42	1	
BR.	L\$114	1	
			5
L\$113:			
LD	T\$12T\$35T\$42	1	s .
			* 1
L\$114: PARM			
PARM	. T\$42	1	* · · · · · · · · · · · · · · · · · · ·
PARM	. T\$26 T\$43	1	• • • • • • • • • • • • • • • • • • •
CALL	GAST\$0	1	
DPARM	3	1	i e
LD	T\$41 T\$0 T\$41	i	· [
		•	15
L\$112:			•••
LD	T\$5T\$41T\$5	1	1
BR	L\$106	1	•
			12 .
L\$105:			
LD	T\$13ABSYT\$13	· 1	• <u></u>
LD SSCAL	S\$3+1T\$14	1	
LD SSCAL BRF	S\$3+1T\$14 .T\$14,L\$117	1	
LD SSCAL BRF SSCAL	S\$3+1T\$14 .T\$14,L\$117 S\$26,.+1,.T\$35	1	H H
LD SSCAL BRF SSCAL LD	\$\$3+1T\$14 .T\$14,L\$117 \$\$26+1T\$35 T\$26T\$35T\$26	1 1 1	H H
LD SSCAL BRF SSCAL	S\$3+1T\$14 .T\$14,L\$117 S\$26,.+1,.T\$35	1	
LD SSCAL BRF SSCAL LD BR	\$\$3+1T\$14 .T\$14,L\$117 \$\$26+1T\$35 T\$26T\$35T\$26	1 1 1	H H
LD SSCAL BRF SSCAL LD BR	S\$3+1T\$14 .T\$14.L\$117 S\$26+1T\$35 T\$26*1\$35T\$26 L\$120	1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1	, , , , , , , , , , , , , , , , , , ,
LD SSCAL BRF SSCAL LD BR	\$\$3+1T\$14 .T\$14,L\$117 \$\$26+1T\$35 T\$26T\$35T\$26	1 1 1	3
LD SSCAL BRF SSCAL LD BR L\$117: LD	S\$3+1T\$14 .T\$14.L\$117 S\$26+1T\$35 T\$26*1\$35T\$26 L\$120	1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1	, , , , , , , , , , , , , , , , , , ,
LD SSCAL BRF SSCAL LD BR L\$117: LD	S\$3+1T\$14 .T\$14.L\$117 S\$26+1T\$35 T\$26*1\$35T\$26 L\$120	1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1	3
LD SSCAL BRF SSCAL LD BR L\$117: LD L\$120: BRF SSCAL	SS31TS14 .TS14.LS117 SS261TS35 TS26TS35TS26 LS120 TS26TS26 .TS26.LS115 SS101TS0	1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1	3
LD SSCAL BRF SSCAL LD BR L\$117: LD L\$120: BRF SSCAL LD	\$5311514 1514.L517 \$526411595 152615351526 L5120 T\$26.01526 1526.L5115 \$51041150	1	3
LD SSCAL BRF SSCAL LD BR L\$117: LD L\$120: BRF SSCAL	SS31TS14 .TS14.LS117 SS261TS35 TS26TS35TS26 LS120 TS26TS26 .TS26.LS115 SS101TS0	1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1	
LD SSCAL BRF SSCAL LD BR L\$117: LD L\$120: BRF SSCAL LD BR BRF SSCAL LD BR	\$5311514 1514.L517 \$526411595 152615351526 L5120 T\$26.01526 1526.L5115 \$51041150	1	
LD SSCAL BRF SSCAL LD BR L\$117: LD L\$120: BRF SSCAL LD BR L\$115:	\$3311514 -1514-1517 -526411525 -15261526 -15261526 -15261526 -15261515 -15261515 -15261516 -15261516 -15261516 -15261516	1	
LD SSCAL BRF SSCAL LD BR L\$117: LD L\$120: BRF SSCAL LD BR L\$115: PARM	\$33411514 -1514.1517 -526411525 -152613.55 -15261.525 -1526.01526 -1526.011150 -15411501541 -1516	1	3
LD SSCAL BRF SSCAL LD BR L\$117: LD L\$120: BRF SSCAL LD BR L\$115: PARM CALL	\$331.1514 \$3261.17514 \$32611.735 \$32611.735 \$326.0.1526 \$326.0.1526 \$326.0.1526 \$326.0.1526 \$326.0.1515 \$31011.730 \$411.730.7541 \$411.730	1	
LD SSCAL BRF SSCAL LD BR L\$117: LD L\$120: BRF SSCAL LD BR L\$115: PARH CALL DPARH	\$33411514 -1514-1517 -526411525 -152613351526 -152601526 -152601526 -152601526 -15261515 -15261515 -15261516 -15261516 -15261516 -15261516 -15261516 -15261516 -15261516 -15261516	1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1	
LD SSCAL BRF SSCAL LD BR L5117: LD L5120: BRF SSCAL LD BR L5115: PARM CALL DPARM PARM	\$331.1514 \$326411525 \$326411525 \$326411525 \$426.01526 \$426.01526 \$426.01526 \$426.041150 \$4411501541 \$4111501541	1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1	3
SSCAL BRF SSCAL LD BR LS117: LD LS120: BRF SSCAL LD BRF SSCAL LD PARM CALL DPARM PARM CALL DPARM CALL DPARM CALL DPARM CALL DPARM CALL DPARM CALL DPARM CALL	\$331.1514 -1514.1517 \$32641.7355 152641.7355 15260.7526 1526.07526 -1526.07526 -1526.45115 -1541750 -1541		
LD SSCAL BRF SSCAL LD BR L5117: LD L5120: BRF SSCAL LD BR L5115: PARM CALL DPARM PARM	\$331.7514 -7514.4517 \$52641.753 \$52641.753 \$1261525 \$1261526 \$1261526 \$1261526 \$1041.750 \$1041.7		
SSCAL BRF SSCAL LD BR SSCAL LD BR SSCAL LD BPF SSCAL LD BPF SSCAL LD PARM CALL DPARM CAL	\$331.1514 -1514.1517 \$32641.7355 152641.7355 15260.7526 1526.07526 -1526.07526 -1526.45115 -1541750 -1541		
LD SSCAL LD BRF SSCAL LD BR L\$1120: BB SSCAL LD BR L\$115: PAPH CALL DPARM PARM CALL DPARM AMD	S331Ts14 S124.Ls17 S1261Ts35 S1261Ts35 S1261Ts26 Ls120 TS26.0.Ts155 S101Ts0 S101Ts0 Ls116		
LSTORE LS	\$331.71514 -71514-1817 -71514-1817 -71514-1817 -7152611.7135 -7152611.7135 -7152611.7135 -7152611.7136 -7152611.7136 -7152611.7136 -7152611.7136 -715261136 -715261136 -715261136 -715261136 -715261136		
LS117: LD LS120: LS117: LD LS120: BP LS117: LD LS120: BP LS115: LD BP SSCAL LD BR LS115: PARH CALL DPARH AND BPF BPF BPF BPF BPF BPF BPF BPF BPF BPF	\$331.71514 \$1514.1517 \$52541.7159 \$52541.7159 \$52541.7159 \$52511.7		
LD LD SSCORE BRF SSCAL LD BR LD	\$5311514 \$52611595 \$52611595 \$52611595 \$52611595 \$52611595 \$5261159 \$526115		
LD SSCAL BRF SSCAL LD BR L\$117: LD L\$120: BR SSCAL LD BR L\$118: BR L\$118: BR L\$115: BR L\$120: BR L\$115: BR L\$120: BR L\$115: BR L\$120: BR	\$331.71514 \$32511.71514 \$1514.1517 \$32511.71525 \$152511.71525 \$152511.71525 \$152511.7152 \$152511.7152 \$11011.7152 \$1		
LD LD SSCORE BRF SSCAL LD BR LD	\$5311514 \$52611595 \$52611595 \$52611595 \$52611595 \$52611595 \$5261159 \$526115		
LS1120: LS120: BF SSCAL LD BR LS117: LD LS120: BF SSCAL LD BP SSCAL LD BP SSCAL LD BP SSCAL LD BP SSCAL LS115: BP PAPH CALL DP-RPH AND BPF BPF LD BPF LD BR LS123: LD	\$331.71514 \$32511.71514 \$1514.1517 \$32511.71525 \$152511.71525 \$152511.71525 \$152511.7152 \$152511.7152 \$11011.7152 \$1		3
LD SSCAL BRF SSCAL LD BR L\$117: LD L\$120: BR SSCAL LD BR L\$118: BR L\$118: BR L\$115: BR L\$120: BR L\$115: BR L\$120: BR L\$115: BR L\$120: BR	\$331.71514 \$32511.71514 \$1514.1517 \$32511.71525 \$152511.71525 \$152511.71525 \$152511.7152 \$152511.7152 \$11011.7152 \$1		

PARM	.x	1	
CALL	GLARTSO	i	
DPARH		1	a .
LD	XT\$0T\$0	1	. •
SSCAL LTSH	S\$15+1T\$12 .T\$12.27T\$12	1	:
OP.	.183318121833	;	:
PARM	-ABSY	i	
CALL	GLTMTSO	1	
DPARH		1	
LD	YTS0,.TS0	1	*
PARM CALL	.150	1	:
DPARM	GHAT\$0	1	:
DR	.1533,.150,.1533	i	i
LD	INSTT\$33T\$33	i	
LD	T\$26YT\$26	1	
BR	L\$122	1	•
L\$121:			25
PARH	.x	1	
CALL	TUMP TSO	i	
DPARH	1	1	
BRF	.T\$0.L\$125	2	
SSCAL	S\$10*1T\$0	1	•
LD BR	T\$33,.1\$0,.T\$33 L\$126	1	
BK	C\$126	1	. 9
L\$125:			
PARM	x	1	
PARM	Y	1	* '
CALL	REGSEARCHTS0	1	
DPARM AND		1	
LD	.X100001T\$12 ABSXT\$12T\$12	2	
AND	.Y100001T\$14	ž	
LD	ABSY T\$14 T\$14	ĭ	
ADD	R.O 1\$35	e	
LD	T\$37,.ABSXT\$37	1	
SSCAL	S\$30+1T\$40	1	•
LD ADD	.T\$35T\$40T\$40	1	
LD	R.1T\$42 T\$37ABSYT\$37	e 1	:
SSCAL	S\$30+1T\$40	i	
LD	.T\$42T\$40T\$40	i	
AND	.T\$40T\$40T\$40	i	9
BRF	. T\$40.L\$127	1	
PARM	-x	1	
PARM CALL	. Y SHOULDEXCHTSO	1	-
DPARM		1	
BRF	.Ts0.L\$131	2	
SSCAL	S\$10+1T\$0	ì	- ·
LD	T\$40T\$0T\$40	1	•
BR	L\$132	1 .	
			. 3
L\$131: SSCAL	5422 .4 7424		
BRE	S\$32+1T\$31 .T\$31.L\$133	1 2	
ADD	X,4,.1845	é	٠.
ADD	RTT\$45T\$45	ĭ	
ADD	.T\$45.0T\$46	ě	
NEQ	T\$46VREGT\$46	1	
AND	.ADDPOSSIBL T\$16 T\$16	1	•
BRF	. T\$46.L\$135	1	
SSCAL LD	S\$1041T\$0	1	
PE.	T\$44,.T\$0,.T\$44 L\$136	1	
J.	C+130	•	
L\$135:			
SSCAL	5\$33.E\$33	1	
S\$33:			
PARM	- ABSX	1	
PARH	- ABSY - ADDPOSSIBL	1	
, restt			

EAL	L GASTSO			
DPA		1		
PAR		1		
CALI	L GNEGTSO	1		
DPAI	RM 1	i		•
RTR	N	i	•	:
E\$33:		•	•	•
LD	T\$44T\$0T\$44	1		
L\$136:				11
LD	T\$43,.T\$44,.T\$43			••
BR	L\$134	1		
	2015.	1		
L\$133:				20
SSCA	L S\$35,E\$35	1		
5\$35:		•		•
BRF LD	ADDPOSSIBL .L\$137	3		
BR	T\$44.270T\$44	1		
DK.	L\$140	1		
L\$137:				2
LD	T\$44,274,.1\$44			
		1		
L\$140:				1
LTSH	.T\$44,33,.T\$44	1		
SSCAL	- S\$12+1T\$16	i		
DR	.T\$44,.T\$16,.T\$44	i		:
RTPN E\$35:		i		:
E#35: PARM				•
CALL	.ABSY GLTRT\$0	1		
DPARM	1 1	1		
PARM	.T\$0	!		
CALL	REGOR. TEO	1		. *
DPARM	1	i		
OR	. T\$44 T\$0 T\$0	i		
LD	INSTTS0TS0	i		:
r.D	T\$43x,.T\$43	1		
L\$134:				20
LD	T\$40,.T\$43,.T\$40			
	1910,.1913,.1910	1		
L\$132:				44
LD	7\$14T\$40,.T\$14	1		
BR	L\$130	i		:
		•		# 55
L\$127: BRF				55
	7\$95.L\$141	3		
SSCAL BRF	S\$32+1T\$31	1		
SSCAL	.T\$31,L\$143 S\$33+1,.T\$0	2		
LD	T\$35T\$0,.7\$35	1		
BR	L\$144	1		
L\$143:				3
SSCAL	5\$35+17\$44	1		
PARM	٠, ٢	1		:
CALL DPARM	MEMOPYATS9	1		
OP-HKH	1 . T\$44,. T\$0 T\$44	1		
LD	INST 1544 1544	1		*
ĹĎ	T\$35X,.T\$35	1		
		1 .		
\$144:				7
LD	T\$40T\$35T\$40	1		
BR	L\$142	i		*
\$141:				15
BRE	Tear Lever			
SSCAL	T\$42.L\$145 \$\$10+1T\$0	3		
LD	T\$35T\$0T\$35	1		
BR .	L\$146	1		
145:				3
SSCAL	S\$32+1T\$31	1		
				•

APPENDIX A

BRF	.T\$31.L\$147	Z			•
BRF	.ADDPOSSIBL.U\$151	3			• .
ESCAL	5\$10+1T\$0	1			:
LD BR	T\$44?50:.1844 L\$152	1			
ъ.	C+10C	•			3
L\$151:					
PARM	•x	. 1			:
CALL	GOLTRTSO	1			-
DPARM LD	1 XTSOTSO	1			- :
SSCAL	S\$3241T\$31	i			
BRF	.T\$31,L\$153	i			
AND	.X100001T\$31	z			•
PARM	.1931	1			
PAPM	ABSY 0	1			
CALL	GASTSO	1			
DPARM	3	i			
PAPM	.T\$0	1			
CALL	GNEGTS0	1			. :
DPARM	1	1			
LD BR	T\$43,.T\$0,.T\$43 L\$154	1			
DK.	C-151	•			12
L\$153:					
PARM	.x	1	-		
PARM	ABSY	1			:
PARM	1	!			:
DPARM	GASTSO 3	1 1			
LD	T\$43,.T\$0T\$43	i			
					. 6
L\$154:					
LD	T\$44T\$43T\$44	1			25
L\$152:				•	
LD	T\$42T\$44T\$42	1			
BP.	L\$150	1			
					33
L\$147: PARM	.x	1			
CALL	PEADYTS0	1			
DPARM	1	i			
BRF	.T\$0.L\$155	2			
BRF	ADDPOSSIBL LSIS?	3			٠.
PAPM CALL	.ABSY GLTRT\$0	1			
DPARM	1	i			
PARM	.150	i			
PARM	.x	1			:
PARM	0	1			
DPARM	GAST\$0 3	1			
LD	TS43T\$9T\$43	i			
BIS	L9160	i			· v
					10
L\$157: PARM	. ABSY	1			
CALI	READY TSO	i			
DPARM	1	i			
BRF	.T\$0.L\$161	ž			
SSCAL	5\$7+1T\$0	1			
PARM	. 150	1			
PARM	. ABSY	1			
CALL	GAST\$0	i			
DPARH	3	i			
LD	T\$31T\$0T\$31	1			
BP	L\$152	1			
1.0151					
L\$161:	. ABSY	1			
CALL	GLTRT\$0	1			
DPARM	1	i			

	AFM .1\$0		
		1	•
		1	
		1	•
	FILL GASTS0	1.	
	PARH 3	1	
	ARM .TSO	1	
E	ALL GNEGTSO	i	
Di	PARM 1	i	7
LE	D T\$31T\$0T\$31	i	
L\$162:			12
L			The second secon
	19131931,1943	1	•
L\$160:			26
- FE			
BR		1	
. 00	P L\$156	1	
			41
L\$155:			
55	CAL 5\$7+1T\$0	1	
	RH .TSO	1	
	RM ABSY	1	
AN		ż	
OR		ī	
PA	RM .T\$43	i	
EA	LL GAST\$0	i	
DP	ern a	i	:
LD		i	
	101111111111111111111111111111111111111		•
L\$156:			19
. LD	**** ****		
	. T\$42 T\$44 T\$42	, 1	
L\$150:			57
LD	T\$35T\$42T\$35	1	
			94
L\$1,46:			-
. LD	T\$10T\$35T\$40	1	
			101
L\$142:			101
· LD	T\$14T\$40T\$14	1	
L\$130:			120
LD	T\$33T\$14T\$33		
	19331.19117.1933	1	•
L\$126>			194
LD	T\$26T\$33T\$26		· · · · · · · · · · · · · · · · · · ·
CU	1926 . 1933 . 1926	1	
L\$122:			263
LD	T\$41T\$26T\$41	1	
			238
L\$116:			
FD	T\$5T\$41T\$5	1	
			250
L\$106:			
LD	T\$30T\$5T\$30	1	
		•	393
L\$100:			303
LD	T\$27,.T\$30,.T\$27		
	19277.19307.1927	1	
L\$60:			323
LD			
LU	T\$11T\$27T\$11	1	
			122
L\$56:			
LD	T\$25,.T\$11,.T\$25	. 1	
			139
L\$50:			
LD	T\$31\$25T\$3	1	
		-	172
L\$16:			
LD	T\$6T\$3T\$6	1	
			178
L\$36: ·		,	
LD	754TS6TS4	1 .	
L\$3Z:		518	l .
	T\$2,.7\$1,.T\$2		
	2197192	1 -	

Please Note:

Pages 135-137, Appendix B: "Algorithm 419, Zeros of a Complex Polynomial [C2]," copyright 1972 by Association for Computing Machinery, Inc., not microfilmed at request of author. Available for consultation at the Carnegie-Mellon University Library.

University Microfilms.

GLOSSARY OF DEFINITIONS

SYMBOL	<u>DEFINED</u> <u>TERM</u>	PAGE
≺	essential ordering	18
∢	initial ordering	20
<	necessary constituent	23
«	essential predecessor	24
•	independent	25
≅	congruence	26
~	strongly similar	28
=	common subexpression	29
β	linear block	31
•	prolog	33
	epilog	33
	postlog	33
۸	formal intersection	36
	interval	49
	cover	50
	redundant expression	57

BIBLIOGRAPHY

- [ASU70] Aho, A.V., Sethi, R. and Ullman, J.D. A formal approach to code optimization, Proc. ACM SIGPLAN Symposium on Compiler Optimization, SIGPLAN Notices, Vol. 5, No. 7 (July 1970), 86-100.
- [AU70] Aho, A.V. and Ullman, J.D. Transformations on straight line programs. Proc. Second Annual ACM Symposium on Theory of Computing, (May 1970), 136-148.
- [AL60] Naur, P. (Ed.) Revised report on the algorithmic language ALGOL 60, Comm. ACM 6, 1 (Jan. 1963), 1-17.
- [A69] Allen, F.E. Program optimization. In <u>Annual Review in Automatic Programming</u>, Vol. 5, Pergamon, New York, 1969.
- [A70] Allen, F.E. Control flow analysis. Proc. ACM SIGPLAN Symposium on Compiler Optimization, SIGPLAN Notices, Vol. 5, No. 7 (July 1970), 1-19.
 - [B71] Bliss reference manual. Computer Science Department Report, Carnegie-Mellon University, Pittsburgh, Pennsylvania, Oct. 1971.
 - [BE69] Busam, V.A. and Englund, D.E. Optimization of expressions in Fortran. Comm. ACM 12, 12 (Dec. 1969), 666-674.
 - [C70] Cocke, J. Global common subexpression elimination. Proc. ACM SIGPLAN Symposium on Compiler Optimization, SIGPLAN Notices, Vol. 5, No. 7 (July 1970), 20-24.
 - [CS70] Cocke, J. and Scinwartz, J.T. Programming Languages and their.

 Compilers, Preliminary Notes, Courant Institute of Mathematical Sciences, New York University, New York, April 1970.

- [D68] Dijkstra, E.W. "Goto statement considered harmful", Letter to the editor, <u>Comm. ACM 11,</u>3 (Mar. 1968).
- [G65] Gear, C.W. High speed complation of efficient object code. <u>Comm. ACM 8</u>, 8 (Aug. 1965), 483-488.
- [GE72] Gerhart, S.L. Verification of APL programs, Ph. D. dissertation, Carnegie-Mellon University, in progress.
- [H72] Hopkins, M.E. A case for the goto, <u>Proc. ACM Annual Conference</u>, Aug. 1972, 787-790.
- [JT72] Jenkins, M.A. and Traub, J.F. Zeros of a complex polynomial, Comm. ACM 15, 2 (Feb. 1972), 97-99.
- [LM69] Lowry, E. and Mediock, C.W. Object code optimization. Comm. ACM 12, 1 (Jan. 1969), 13-22.
- [N65] Nievergelt, J. On the automatic simplification of computer programs. Comm. ACM 8, 6 (June 1965), 366-370.
- [P70] PDP10 reference manual. Digital Equipment Corporation, Maynard, Massachusetts, 1970.
- [SS69] Shapiro, R.M. and Saint, H. The representation of algorithms.

 | Rome | Air | Development | Center | Technical | Report, RADC-TR-69-313, Vol. 2, Sept. 1969.
- [W71] Wulf, W.A. Programming without the goto, IFIP, I971.
- [W72] Wulf, W.A. A case against the goto, Proc. ACM Annual Conference, Aug. 1972, 791-796.
- [WRH71] Wulf, W.A., Russell, D.B. and Habermann, A.N. Bliss: a language for systems programming. Comm. ACM 14, 12 (Dec. 1971), 780-790.